

DATA SCIENCE FRAMEWORKS AND MANAGED SERVICES: WHEN TO AVOID THE SHINY NEW TOYS



JON TUTCHER, BBC DATALAB, PYDATA LONDON

14 JULY 2019

On DAG: pydata_london

schedule: None

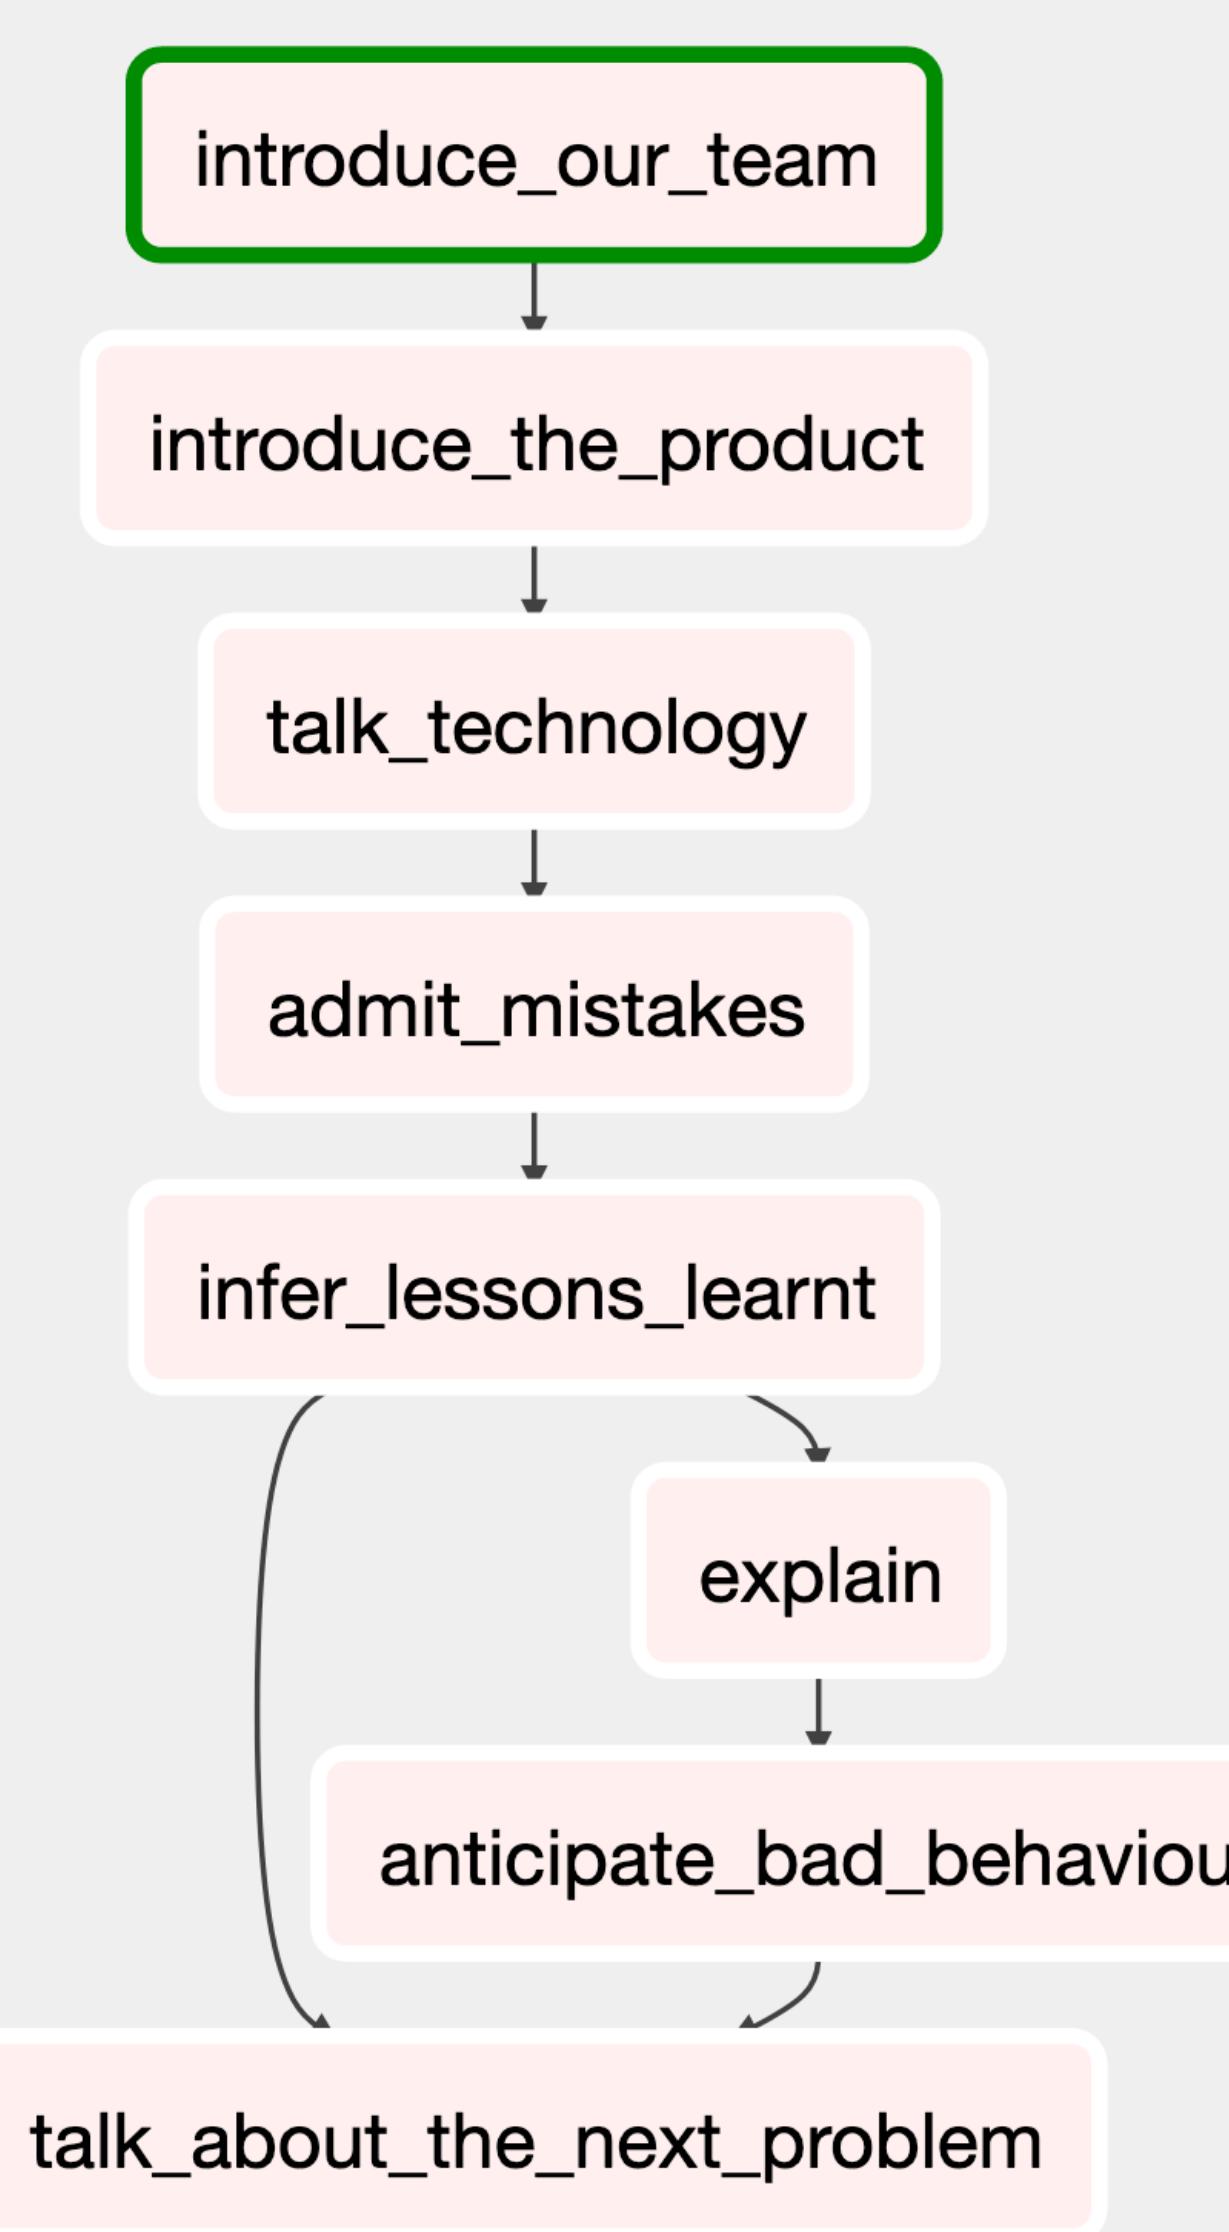
[Graph View](#) [Tree View](#) [Task Duration](#) [Task Tries](#) [Landing Times](#) [Gantt](#) [Details](#) [Code](#) [Trigger DAG](#) [Refresh](#) [Delete](#)

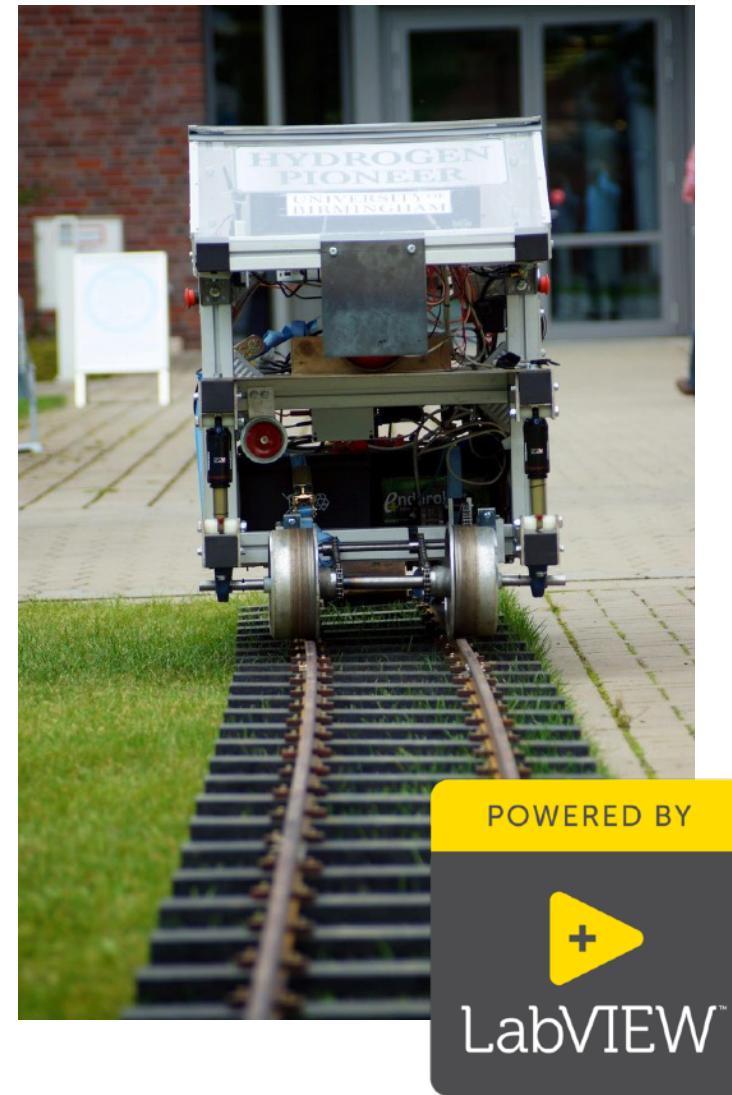
running Base date: 2019-07-13 18:38:04 Number of runs: 25 Run: manual_2019-07-13T18:38:03.302273+00:00 Layout: Top->Bottom Go

Search for...

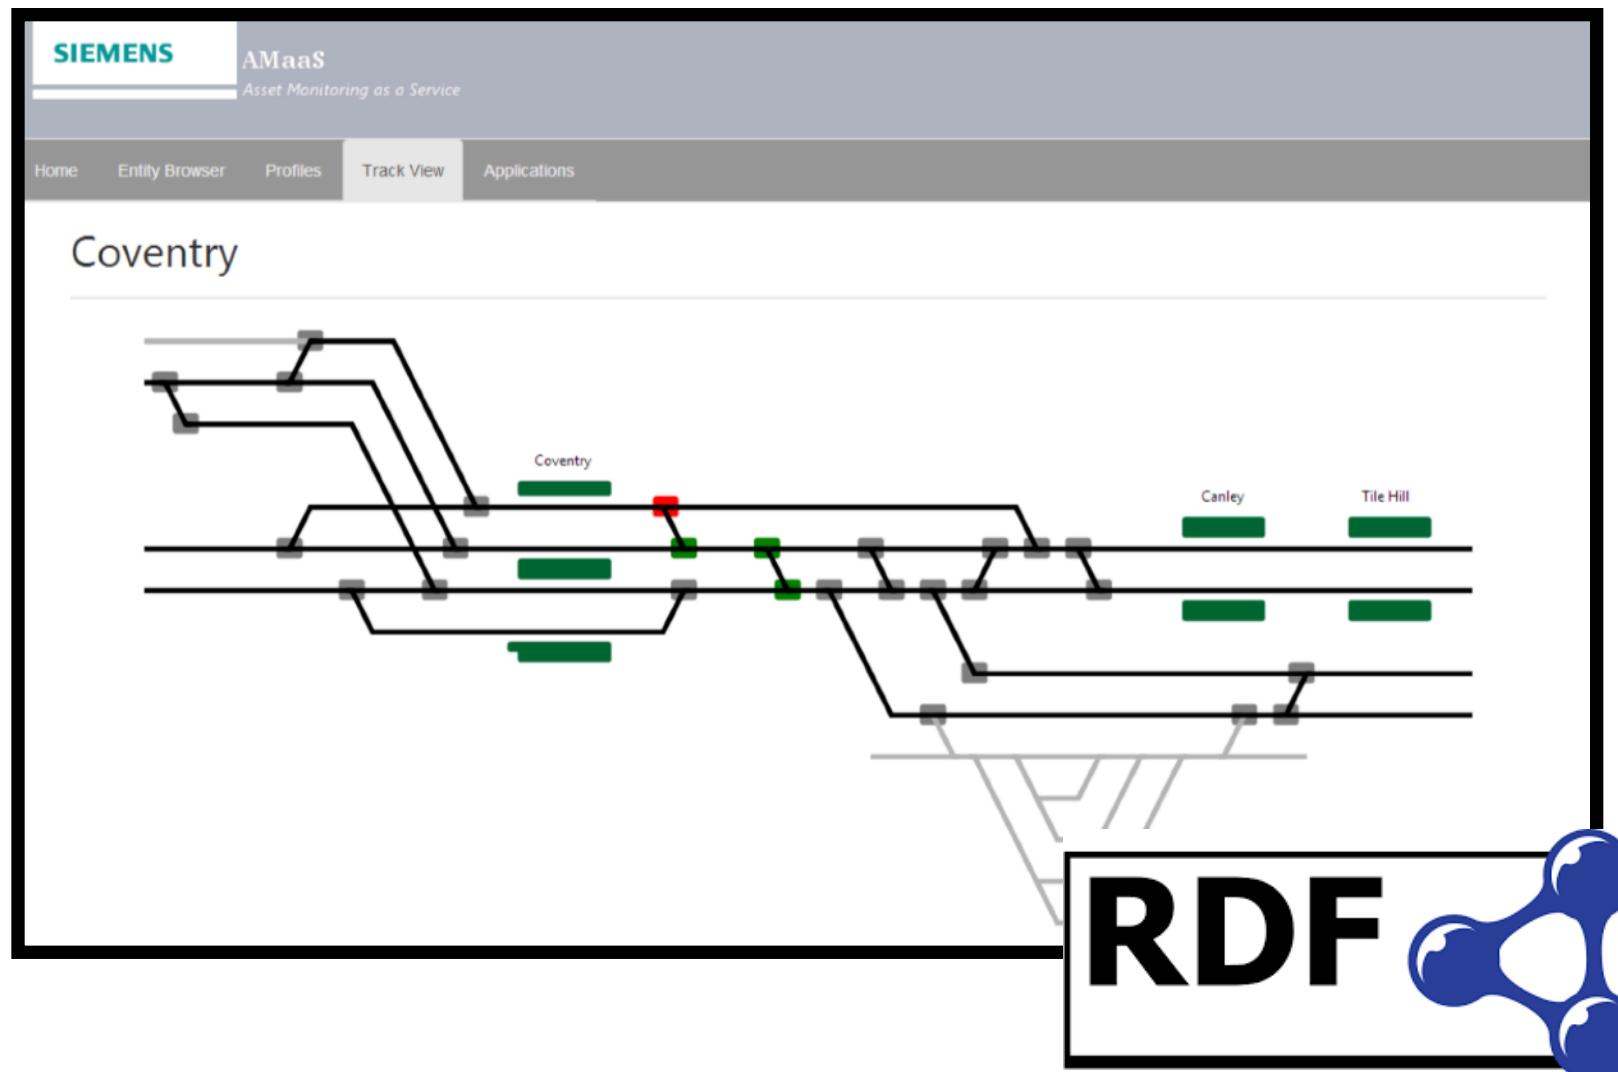
PythonOperator

(success) (running) (failed) (skipped) (up_for_reschedule) (up_for_retry) (queued) no_status

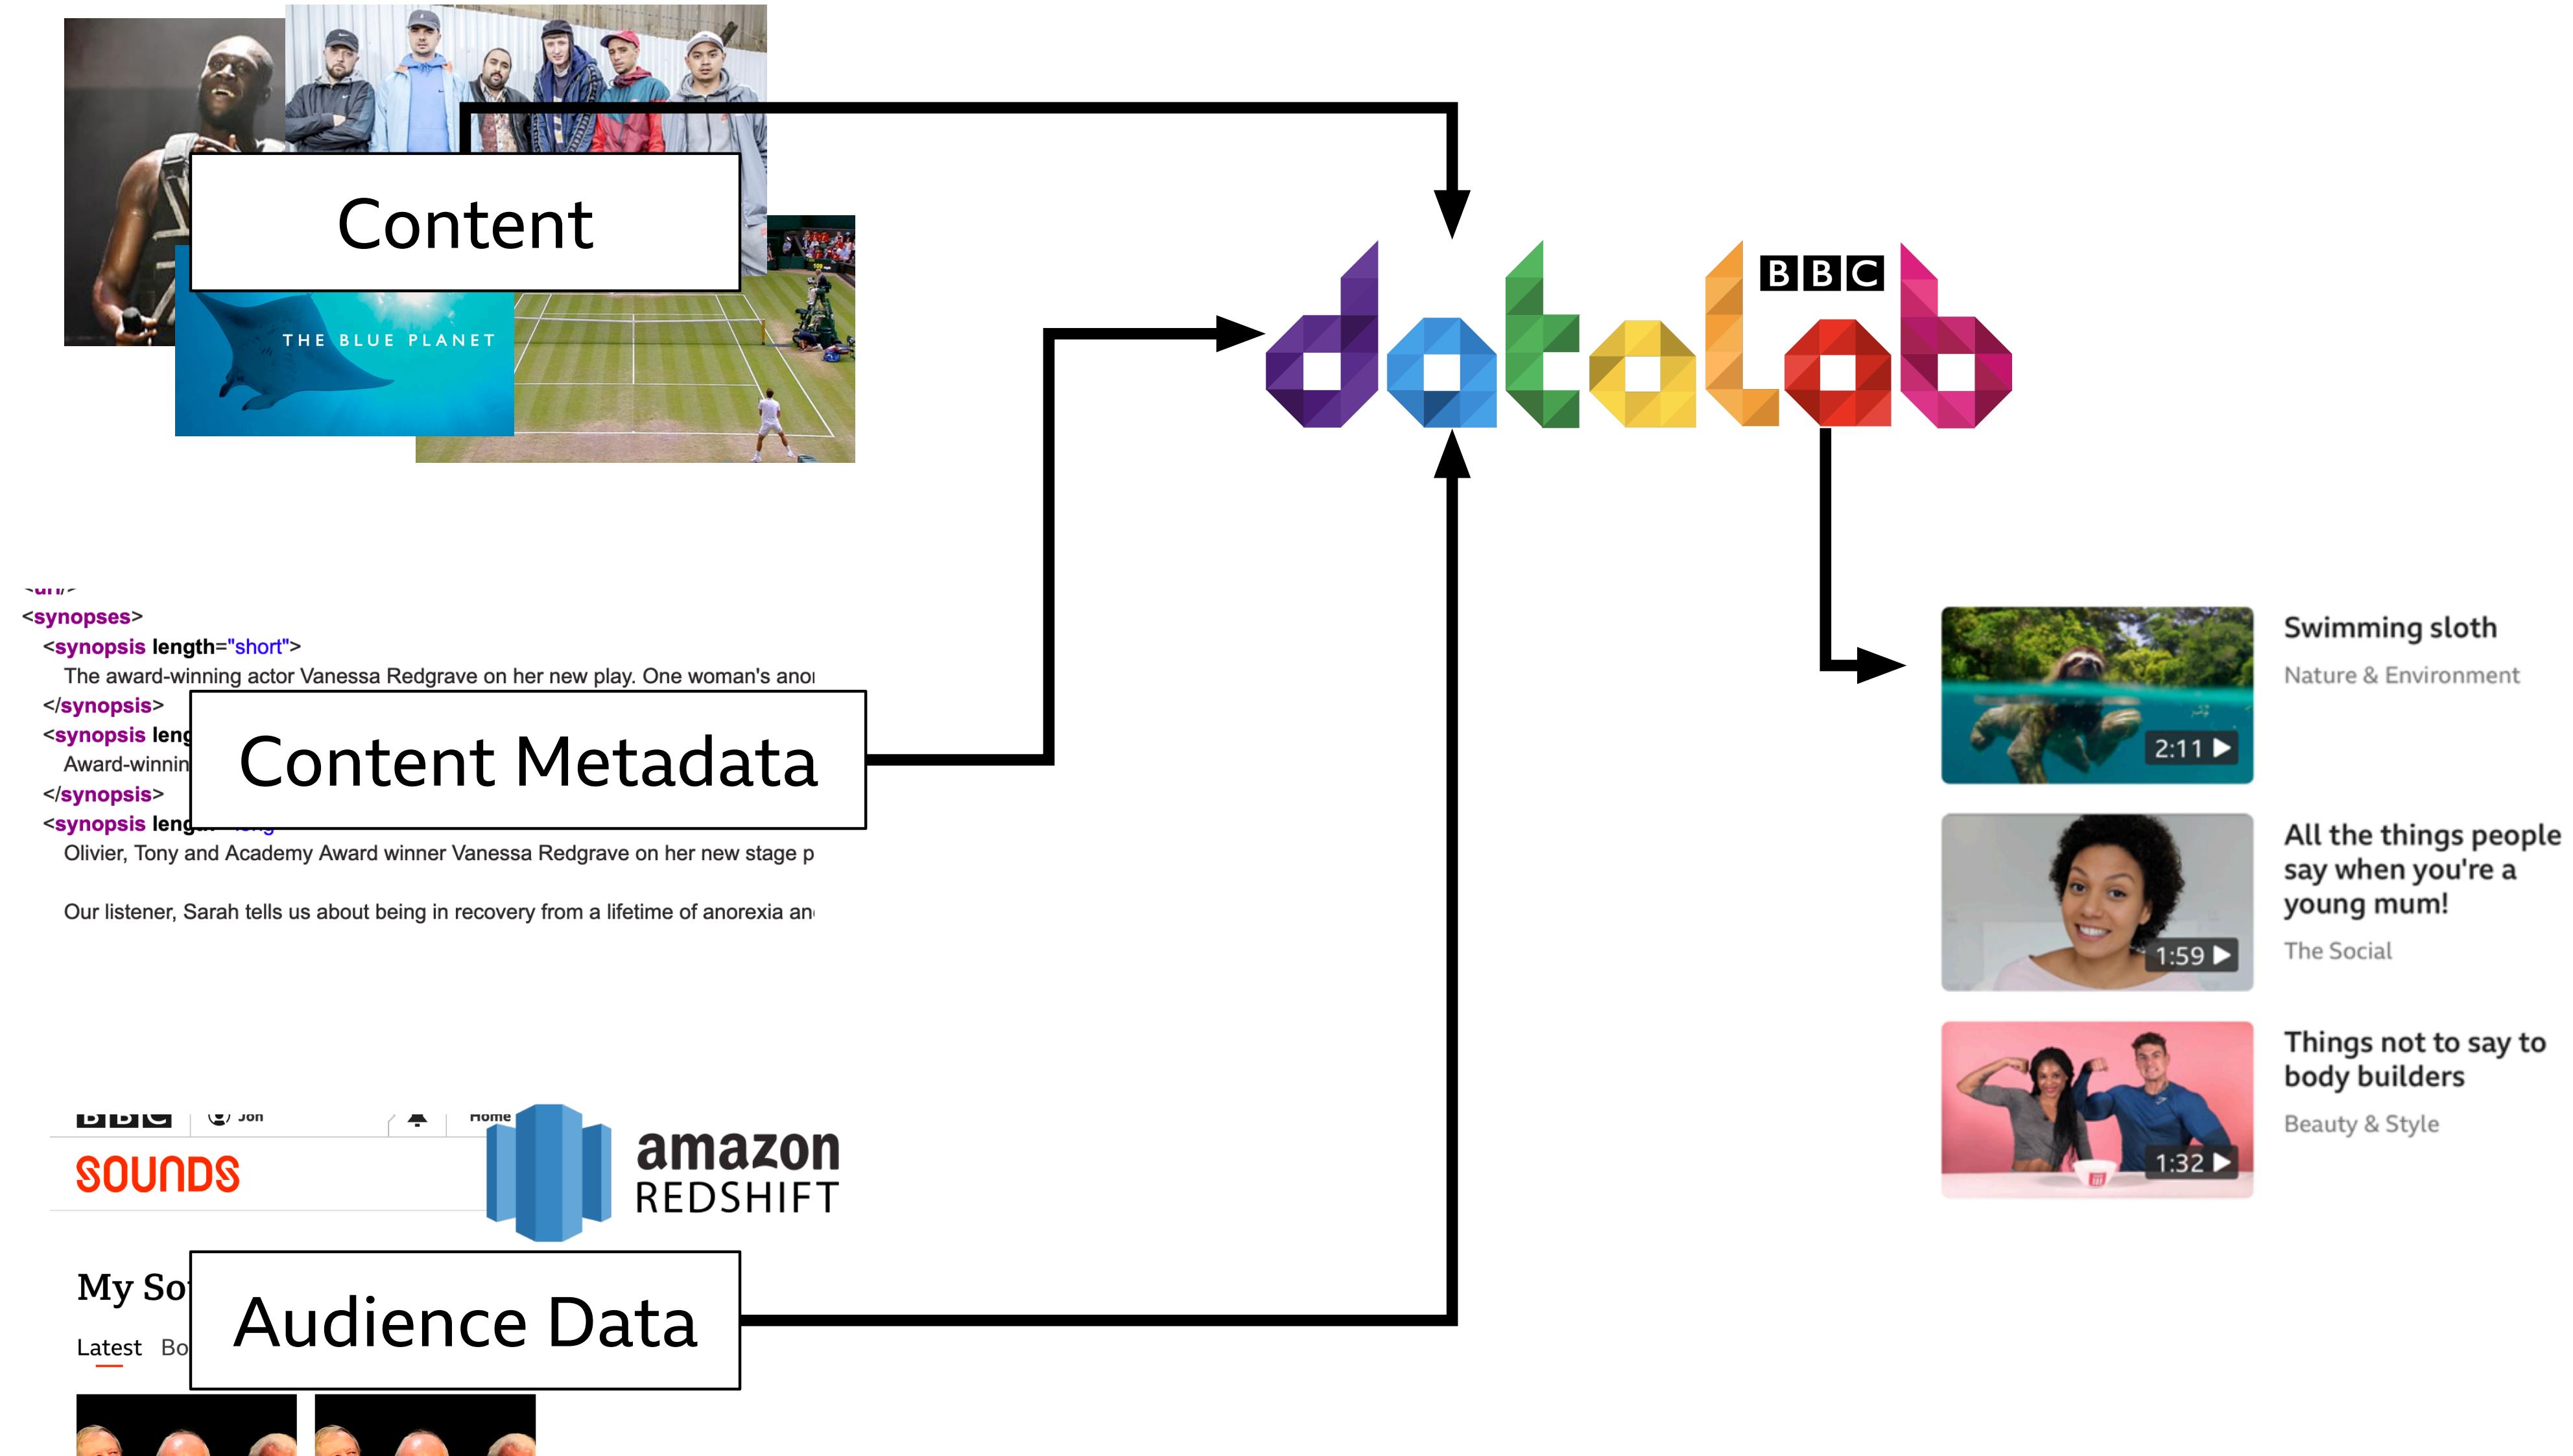




 **GNURadio**
THE FREE & OPEN SOFTWARE RADIO ECOSYSTEM







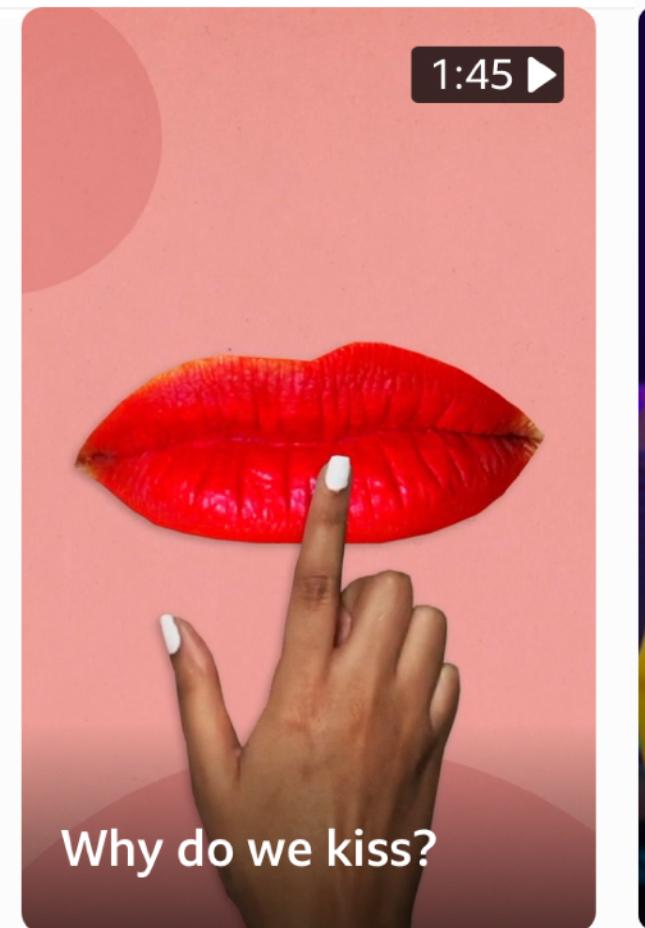
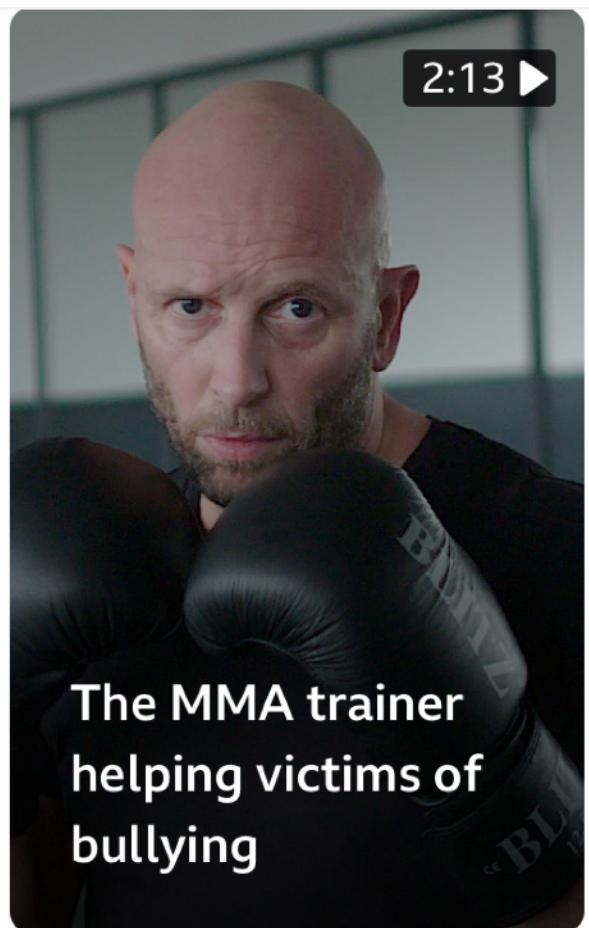
9:41



TUE 7 MAY

Good Morning

Videos of the day



Money



The city that gives you free beer for cycling

42 minutes ago

BBC World Hacks



Related



How to eat carbs and stay healthy

Food & Drink



Five compelling reasons why we all need to sleep more

Health & Wellbeing



The malicious influence of the Spice Girls

Stand-up

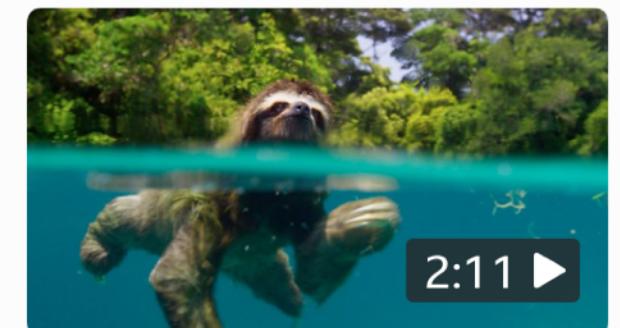
9:41



Inspiring

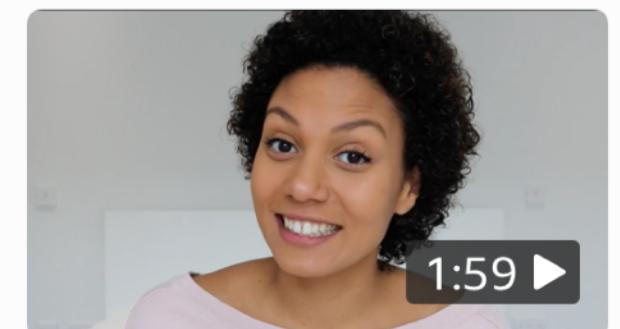


Cancel



Swimming sloth

Nature & Environment



All the things people say when you're a young mum!

The Social



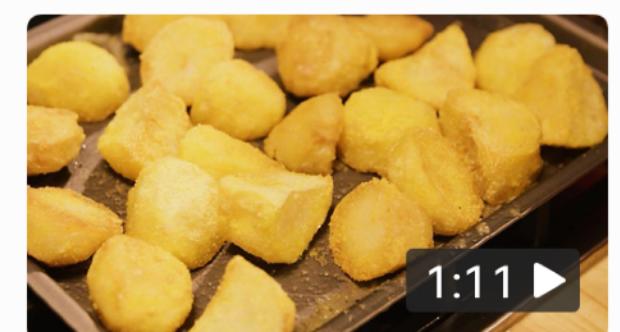
Things not to say to body builders

Beauty & Style



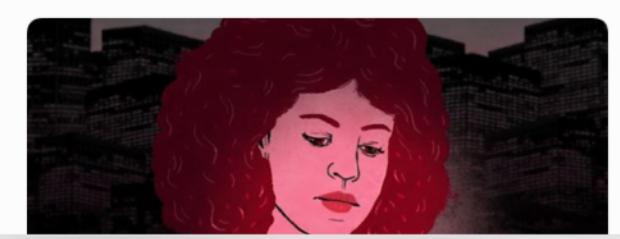
Tips for success by the youngest British woman to climb Everest

Top Tips

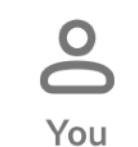


How to make extra crunchy roast potatoes

Food



Ideas



BBC+

GETTING STARTED

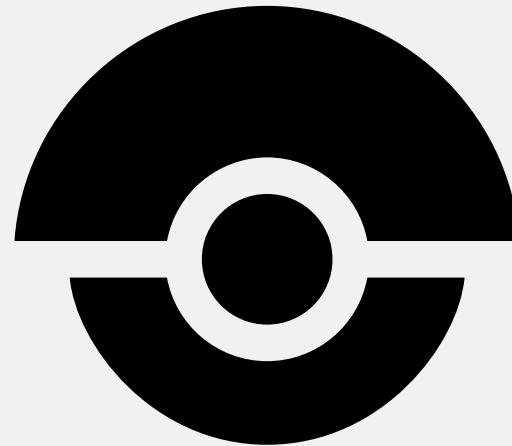
"We'll have other customers - we need to build a platform"

"We're a new, independent team - the bosses want us to try some new technology"

"We don't really know what the requirements are yet, so let's build something really flexible"

BBC+

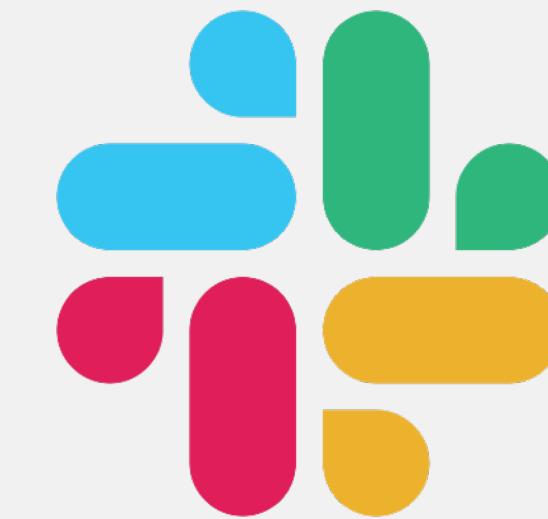
OUR RESPONSE



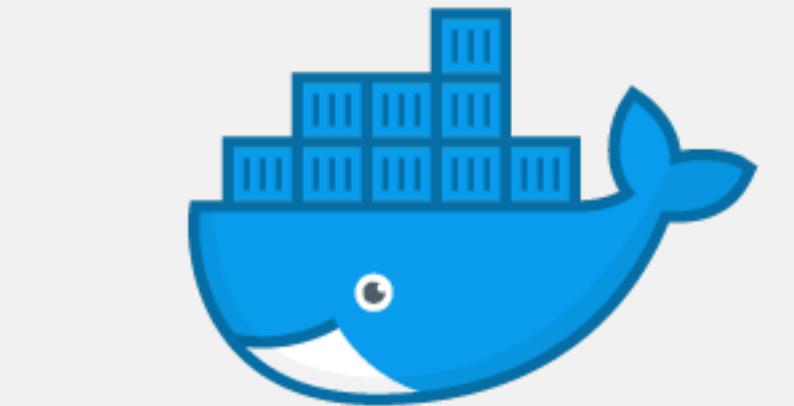
DroneCI



kubernetes

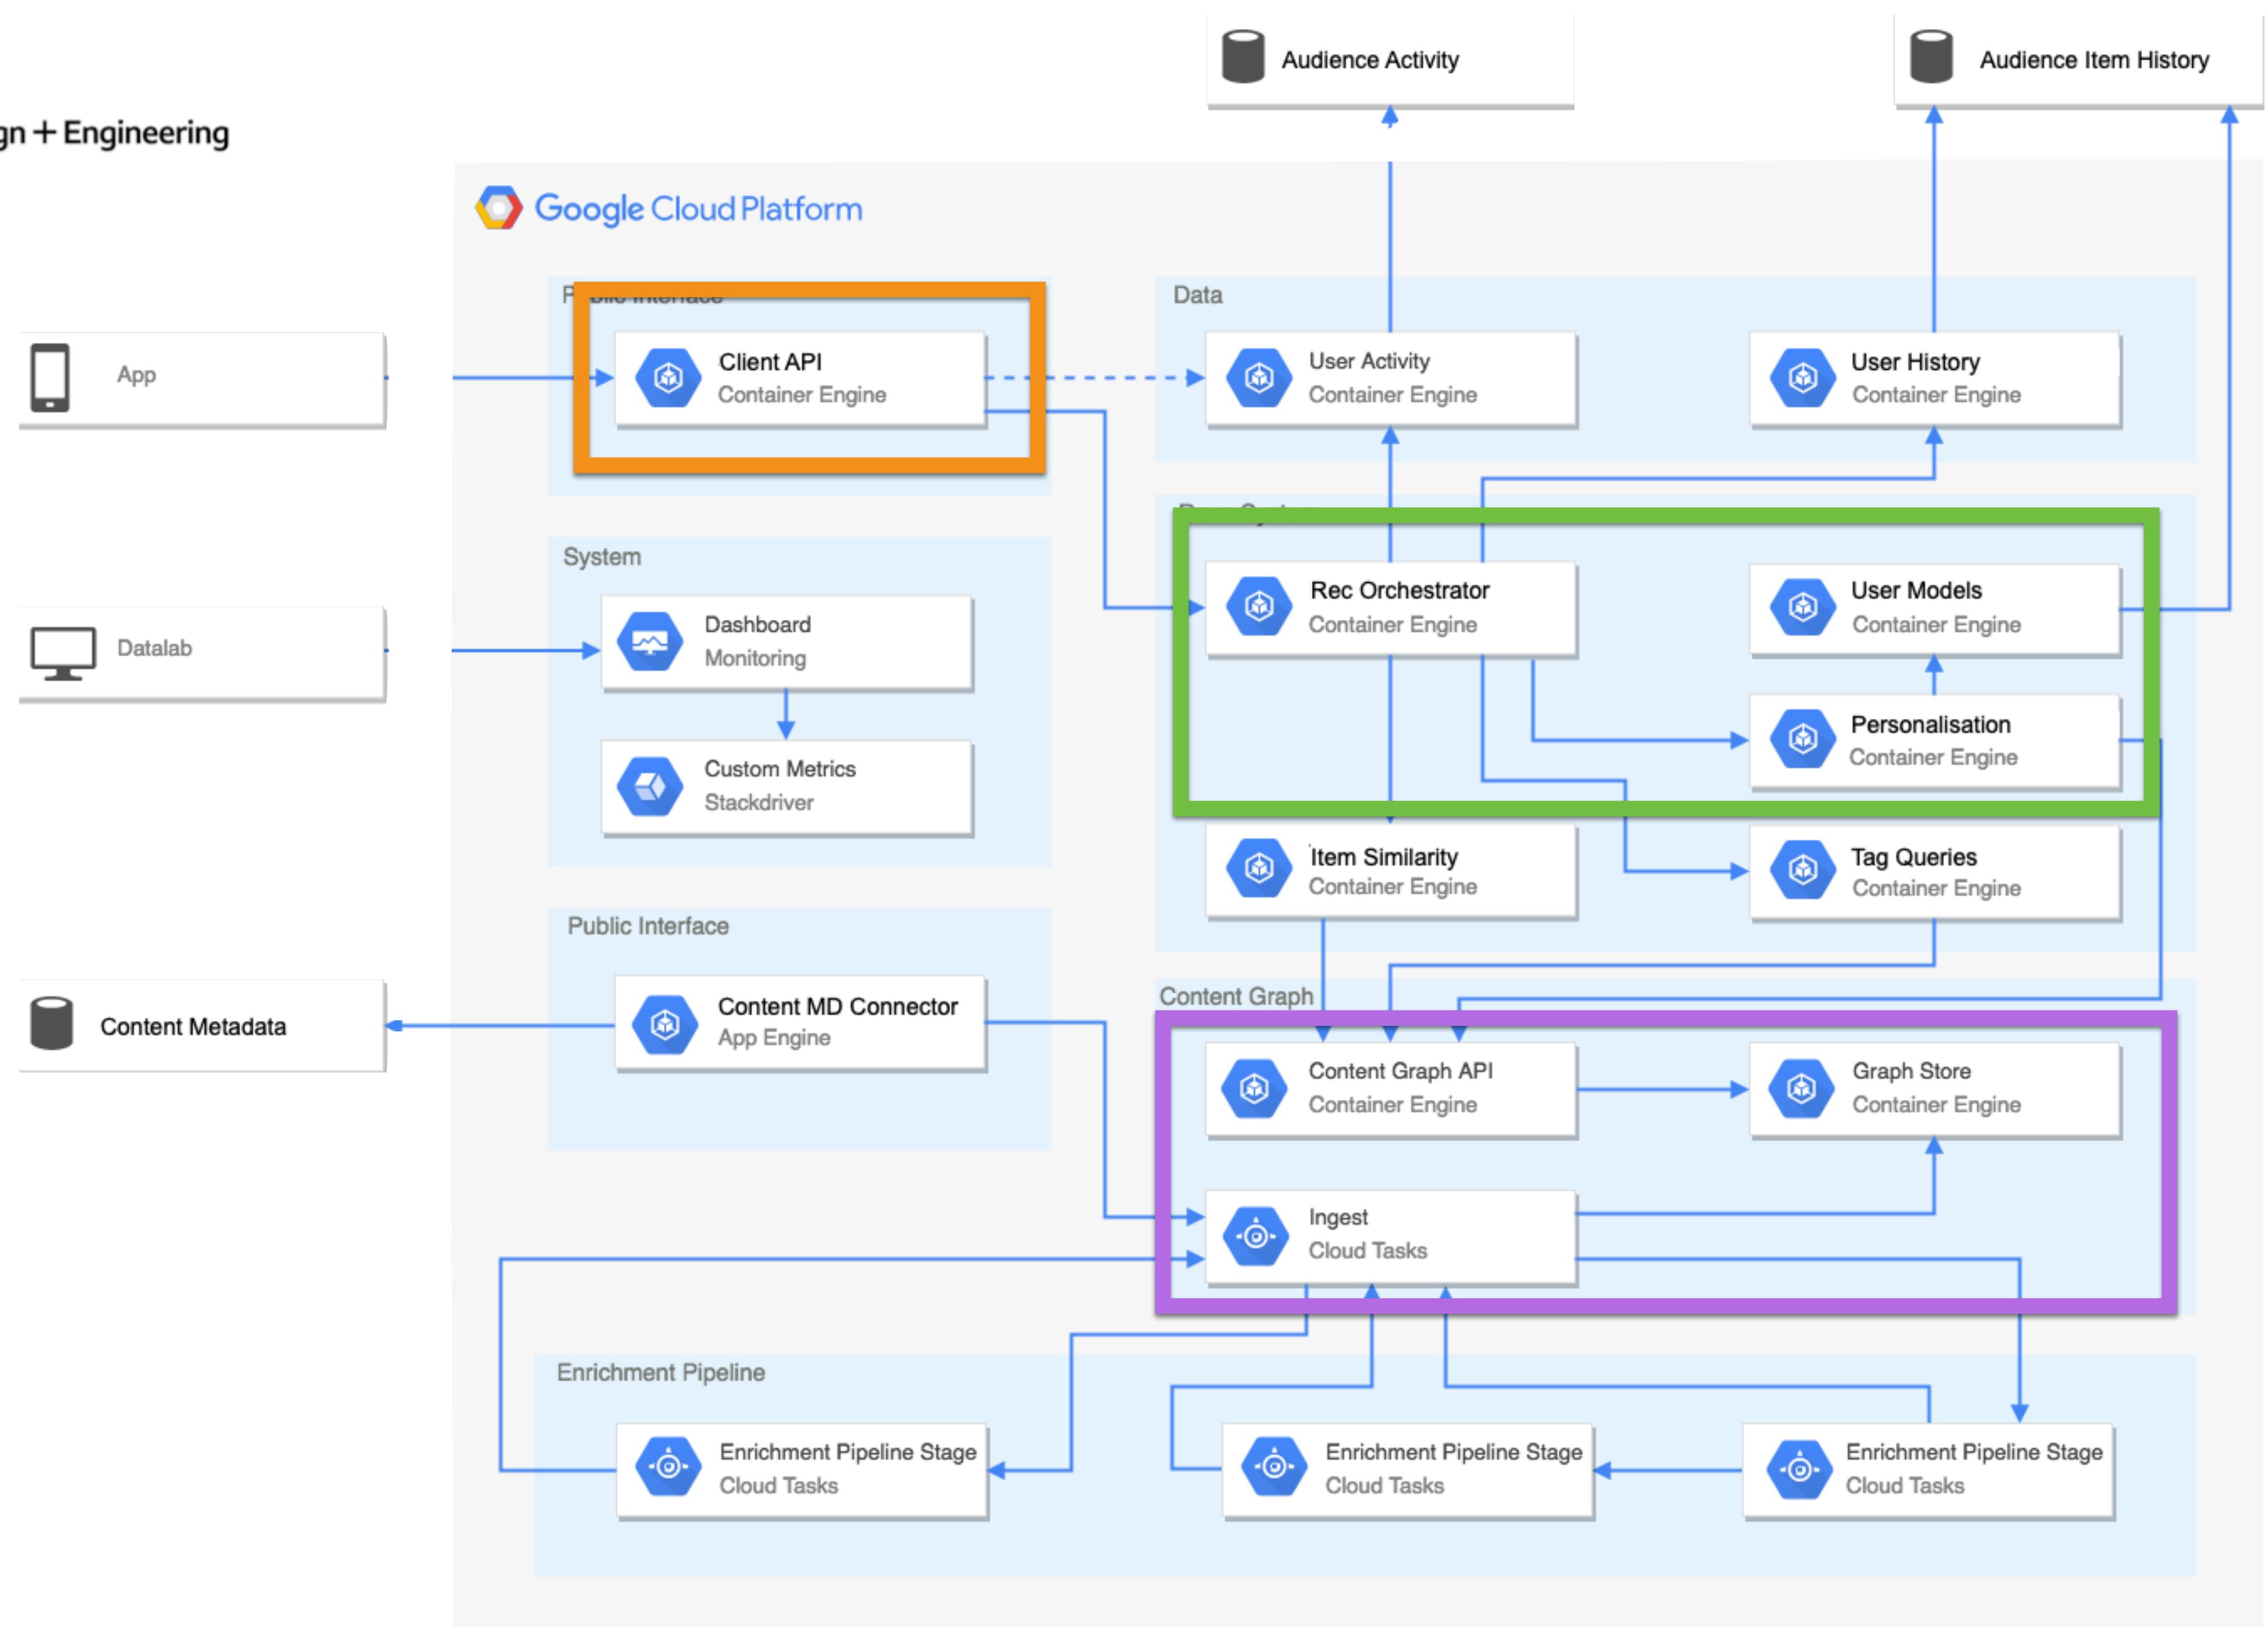


Google Cloud



docker

OUR TECHNOLOGY CHOICES





Search or jump to...

/

Pull requests Issues Marketplace Explore



+



BBC / TS&A / Datalab

Discussions

Members 24

Teams 3

Repositories 80

Projects 0

Settings

Find a repository...

Add repository

Select all

[bbc/connected-data](#) Private updated 3 days ago



Admin ▾

[bbc/connected-data-acceptance-tests](#) Private
updated on Apr 25



Admin ▾

[bbc/connected-data-amanita](#) Private updated on May
28



Admin ▾

[bbc/connected-data-bramble](#) Private updated on May
28



Admin ▾

[bbc/connected-data-bristlecone](#) Private updated on
May 28



Admin ▾

[bbc/connected-data-catopsis](#) Private updated on May



Admin ▾

```
1 # Generated by the gRPC Python protocol compiler plugin. DO NOT EDIT!
2 import grpc
3
4 import app.bramble_pb2 as bramble_pb2
5
6
7 class BrambleServiceStub(object):
8     pass
9
10    def __init__(self, channel):
11        """Constructor.
12
13        Args:
14            channel: A grpc.Channel.
15        """
16
17        self.ListRecommendations = channel.unary_unary(
18            '/bramble.BrambleService/ListRecommendations',
19            request_serializer=bramble_pb2.ListRecommendationsRequest.SerializeToString,
20            response_deserializer=bramble_pb2.Recommendations.FromString,
21        )
22        self.HealthCheck = channel.unary_unary(
23            '/bramble.BrambleService/HealthCheck',
24            request_serializer=bramble_pb2.Empty.SerializeToString,
25            response_deserializer=bramble_pb2.Empty.FromString,
26        )
27        self.ListUserHistory = channel.unary_unary(
28            '/bramble.BrambleService/ListUserHistory',
```

```
clip_recommendations = bramble.ListRecommendations(user_id, media_type, time)
```

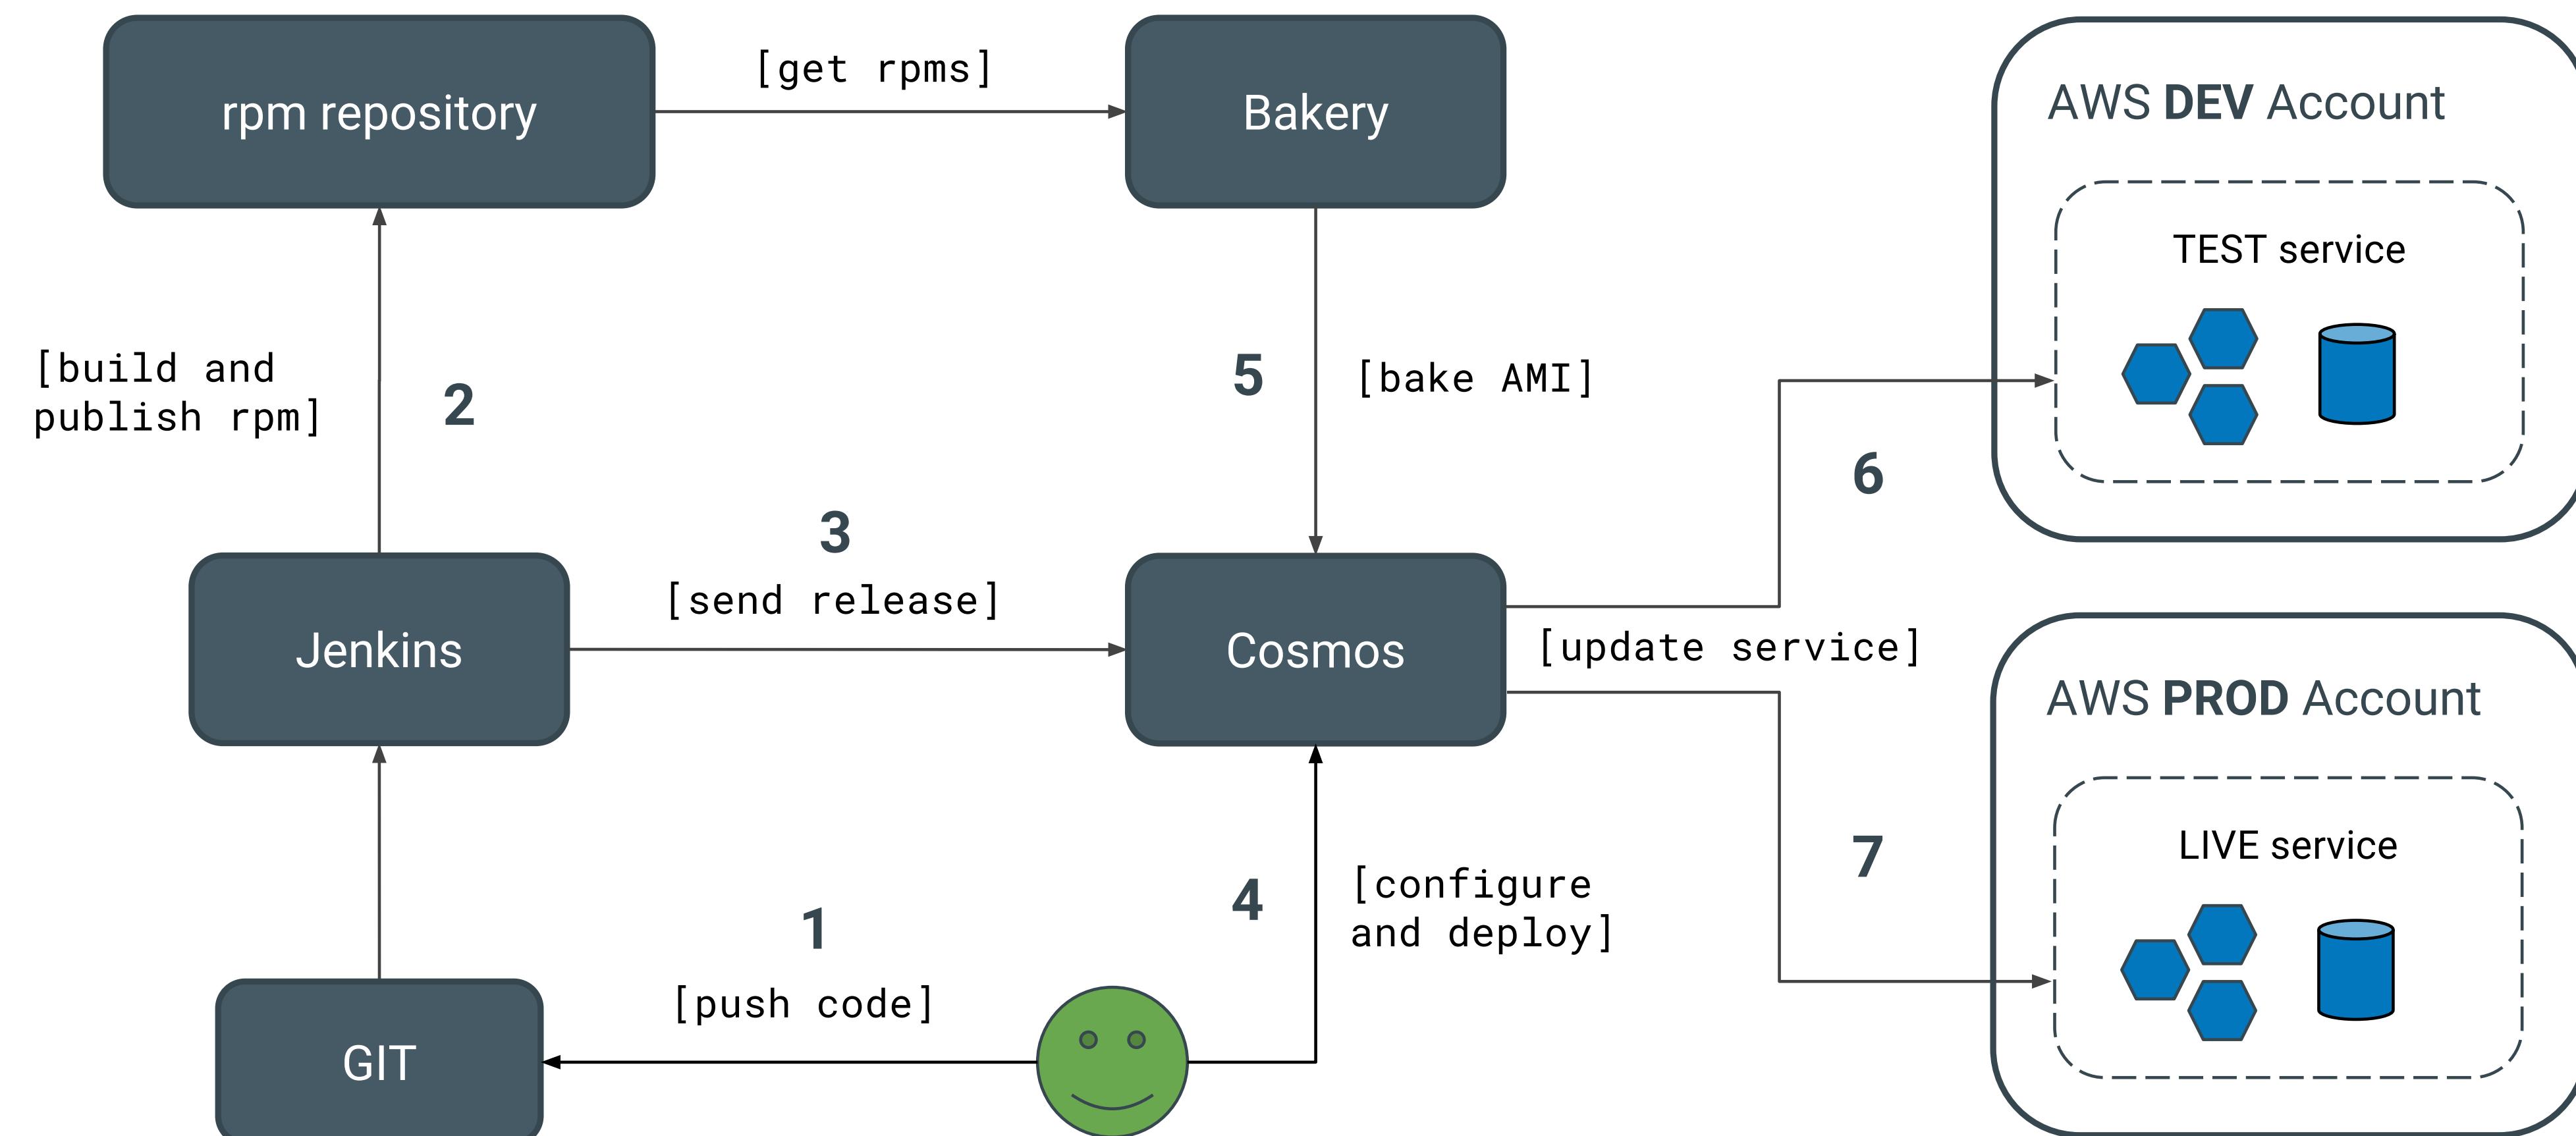
TECHNOLOGY CHOICES

GRPC

HTTP (JSON)	gRPC (protobuf)
Every developer knows how to use (ish)	Developers need training
Tooling is everywhere	Tooling is difficult
Loads of python libraries!	gRPC library (un-googleable)
Slow?	Fast Slow (in python)*
API changes are tricky	API changes backward-compatible

TECHNOLOGY CHOICES

BBC "TRADITIONAL" SOFTWARE DEPLOYMENT



GOOGLE CLOUD PLATFORM

AI and machine learning

[Text-to-Speech](#) · [Vision](#) · [Translation](#) · More

Hybrid and multi-cloud

[Anthos](#) · [GKE On-Prem](#) · [Istio on GKE](#) · More

API management

[Apigee API Platform](#) · [Cloud Endpoints](#) · More

Internet of Things

[Cloud IoT Core](#) · [Edge TPU](#)

Compute

[Compute Engine](#) · [App Engine](#) · More

Migration

[Data Transfer](#) · [Transfer Appliance](#) · More

Data analytics

[BigQuery](#) · [Cloud Datalab](#) · More

Networking

[DNS](#) · [CDN](#) · [Virtual Private Cloud](#) · More

Databases

[Cloud SQL](#) · [Cloud Datastore](#) · More

Security

[Security Key Enforcement](#) · [Cloud IAM](#) · More

Developer Tools

[Container Registry](#) · [Cloud SDK](#) · More

Storage

[Cloud Storage](#) · [Persistent Disk](#) · More

MORE CLOUD PRODUCTS

G Suite

Gmail, Docs, Drive, Hangouts, and more

Google Maps Platform

Build with real-time, comprehensive data

Cloud Identity

Easily manage user identities

Chrome Enterprise

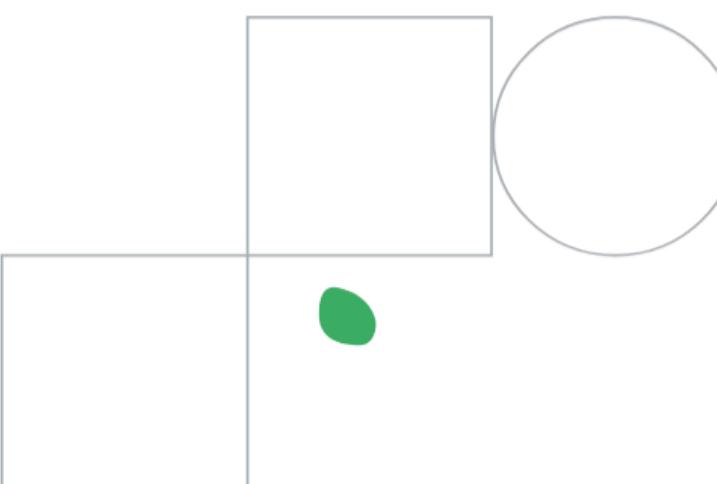
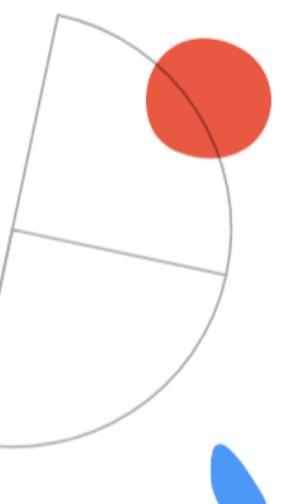
Get Chrome OS devices and browser

Android Enterprise

Intelligent devices, OS, and business apps

Hire by Google

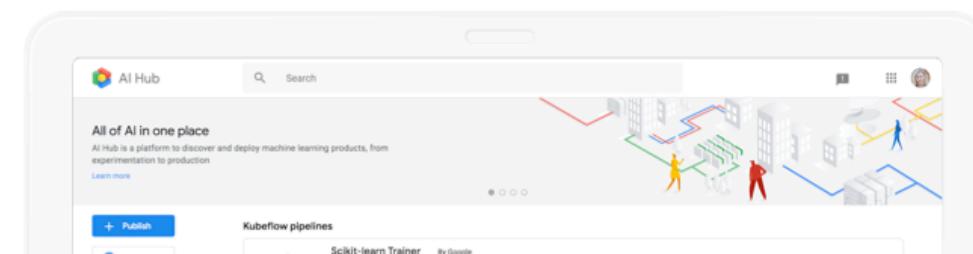
Identify, evaluate, and hire better

[See all products \(100+\)](#)

AI Hub

Hosted repository of plug-and-play AI components

Google Cloud's AI Hub provides enterprise-grade sharing capabilities, including end-to-end AI pipelines and out-of-the-box algorithms, that let your organization privately

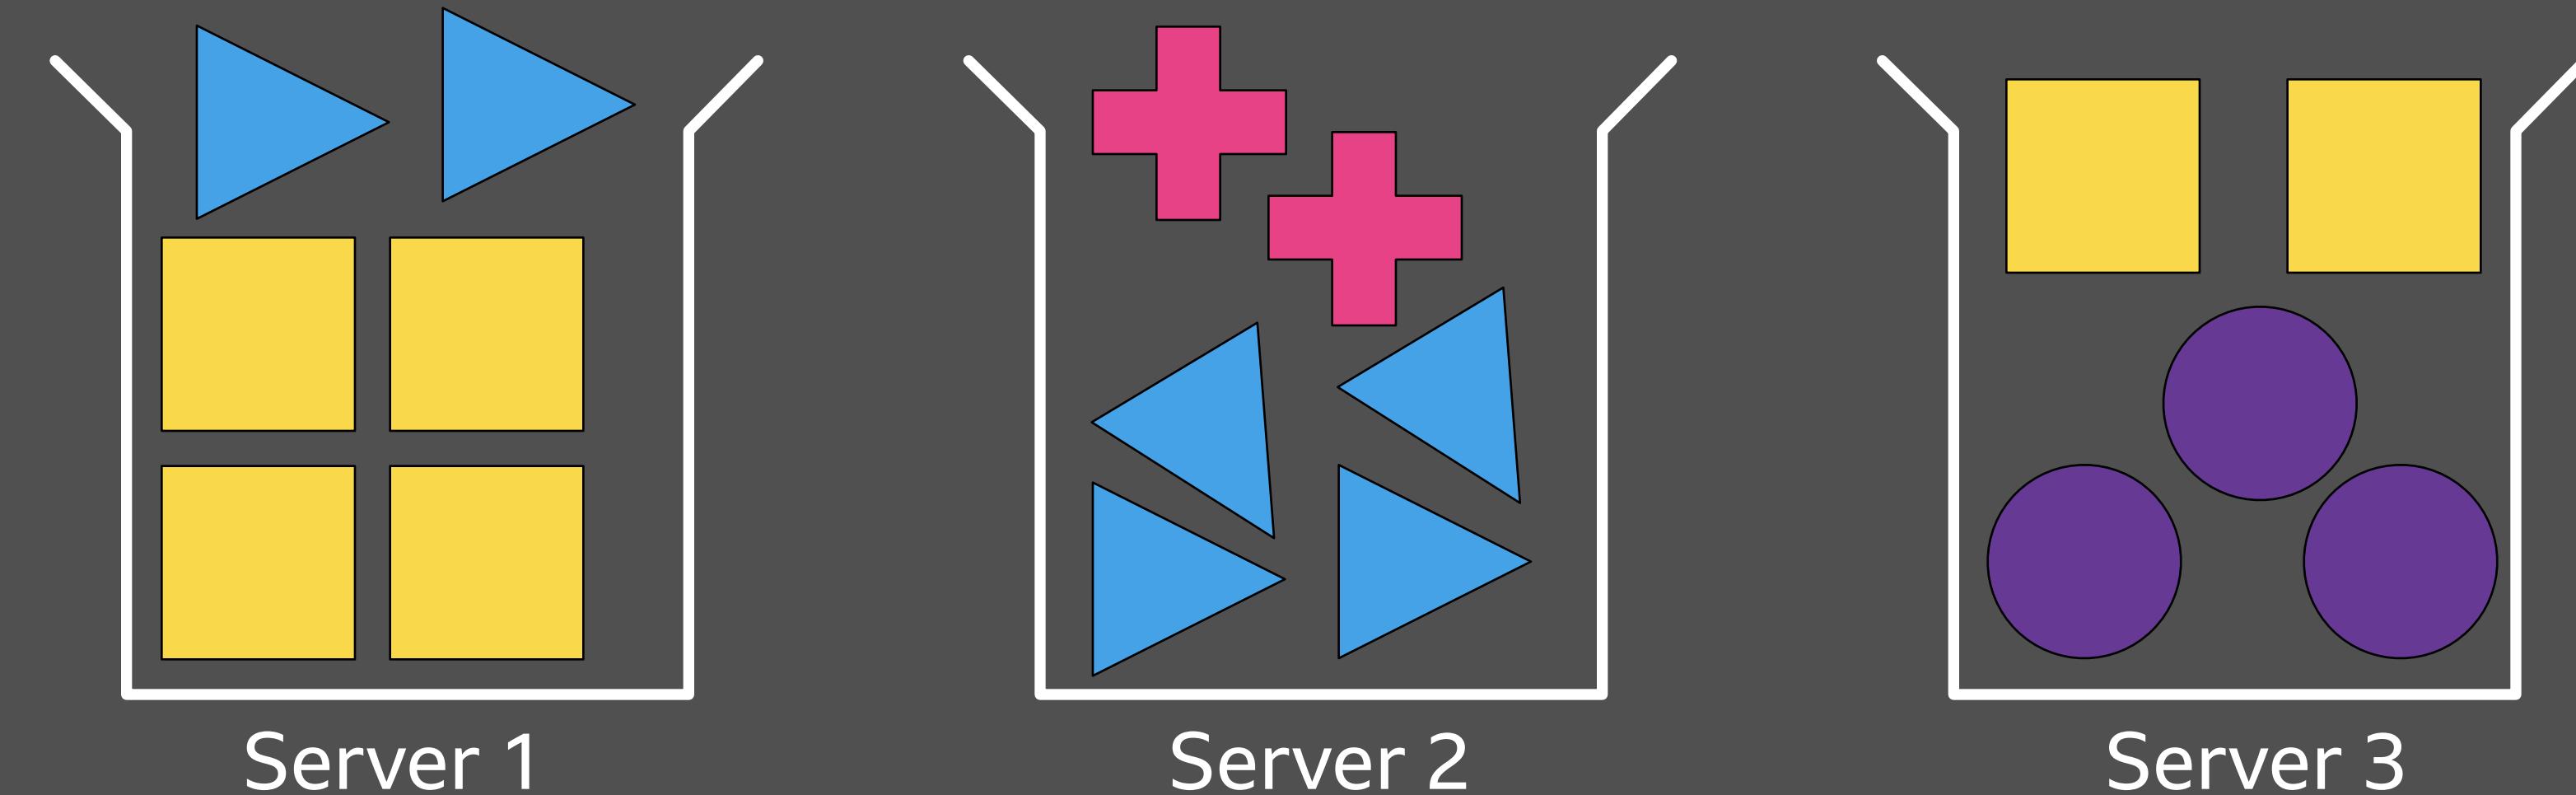


TECHNOLOGY CHOICES

DOCKER & KUBERNETES

```
1 FROM microservice-base:latest
2 MAINTAINER BBC Datalab <datalab@bbc.co.uk>
3
4 RUN pip3 install --upgrade pipenv && pipenv install --deploy --system --verbose
5
6 CMD /usr/local/bin/gunicorn --bind 0.0.0.0:5000 --workers 2 --access-logfile - app:server
```

some_service ➔ ↗ master • ? ↑1 ➔ kubectl apply -f service.yml



datalabservice

PIPELINES

CLUSTERS

TASKS

LOAD BALANCERS

SECURITY GROUPS

CONFIG

Clusters

[Edit multiple](#)[Show](#) Instances with details[Create Server Group](#)

MY-K8S-ACCOUNT

datalabservice-production

1 ▲ / 1 - : 50%

DEFAULT

V006: cg-pov/datalab-service:1478ea85412204708419bb99f419e466584938c0

1 ▲ : 100%

V005: cg-pov/datalab-service:545246b8a203e6dbaa632033450923b84005dd91

1 - : 0%

MY-K8S-ACCOUNT

datalabservice-stage

1 ▲ : 100%

DEFAULT

V008: cg-pov/datalab-service:1478ea85412204708419bb99f419e466584938c0

1 ▲ : 100%

datalabservice-live

[Load Balancer Actions ▾](#)

LOAD BALANCER DETAILS

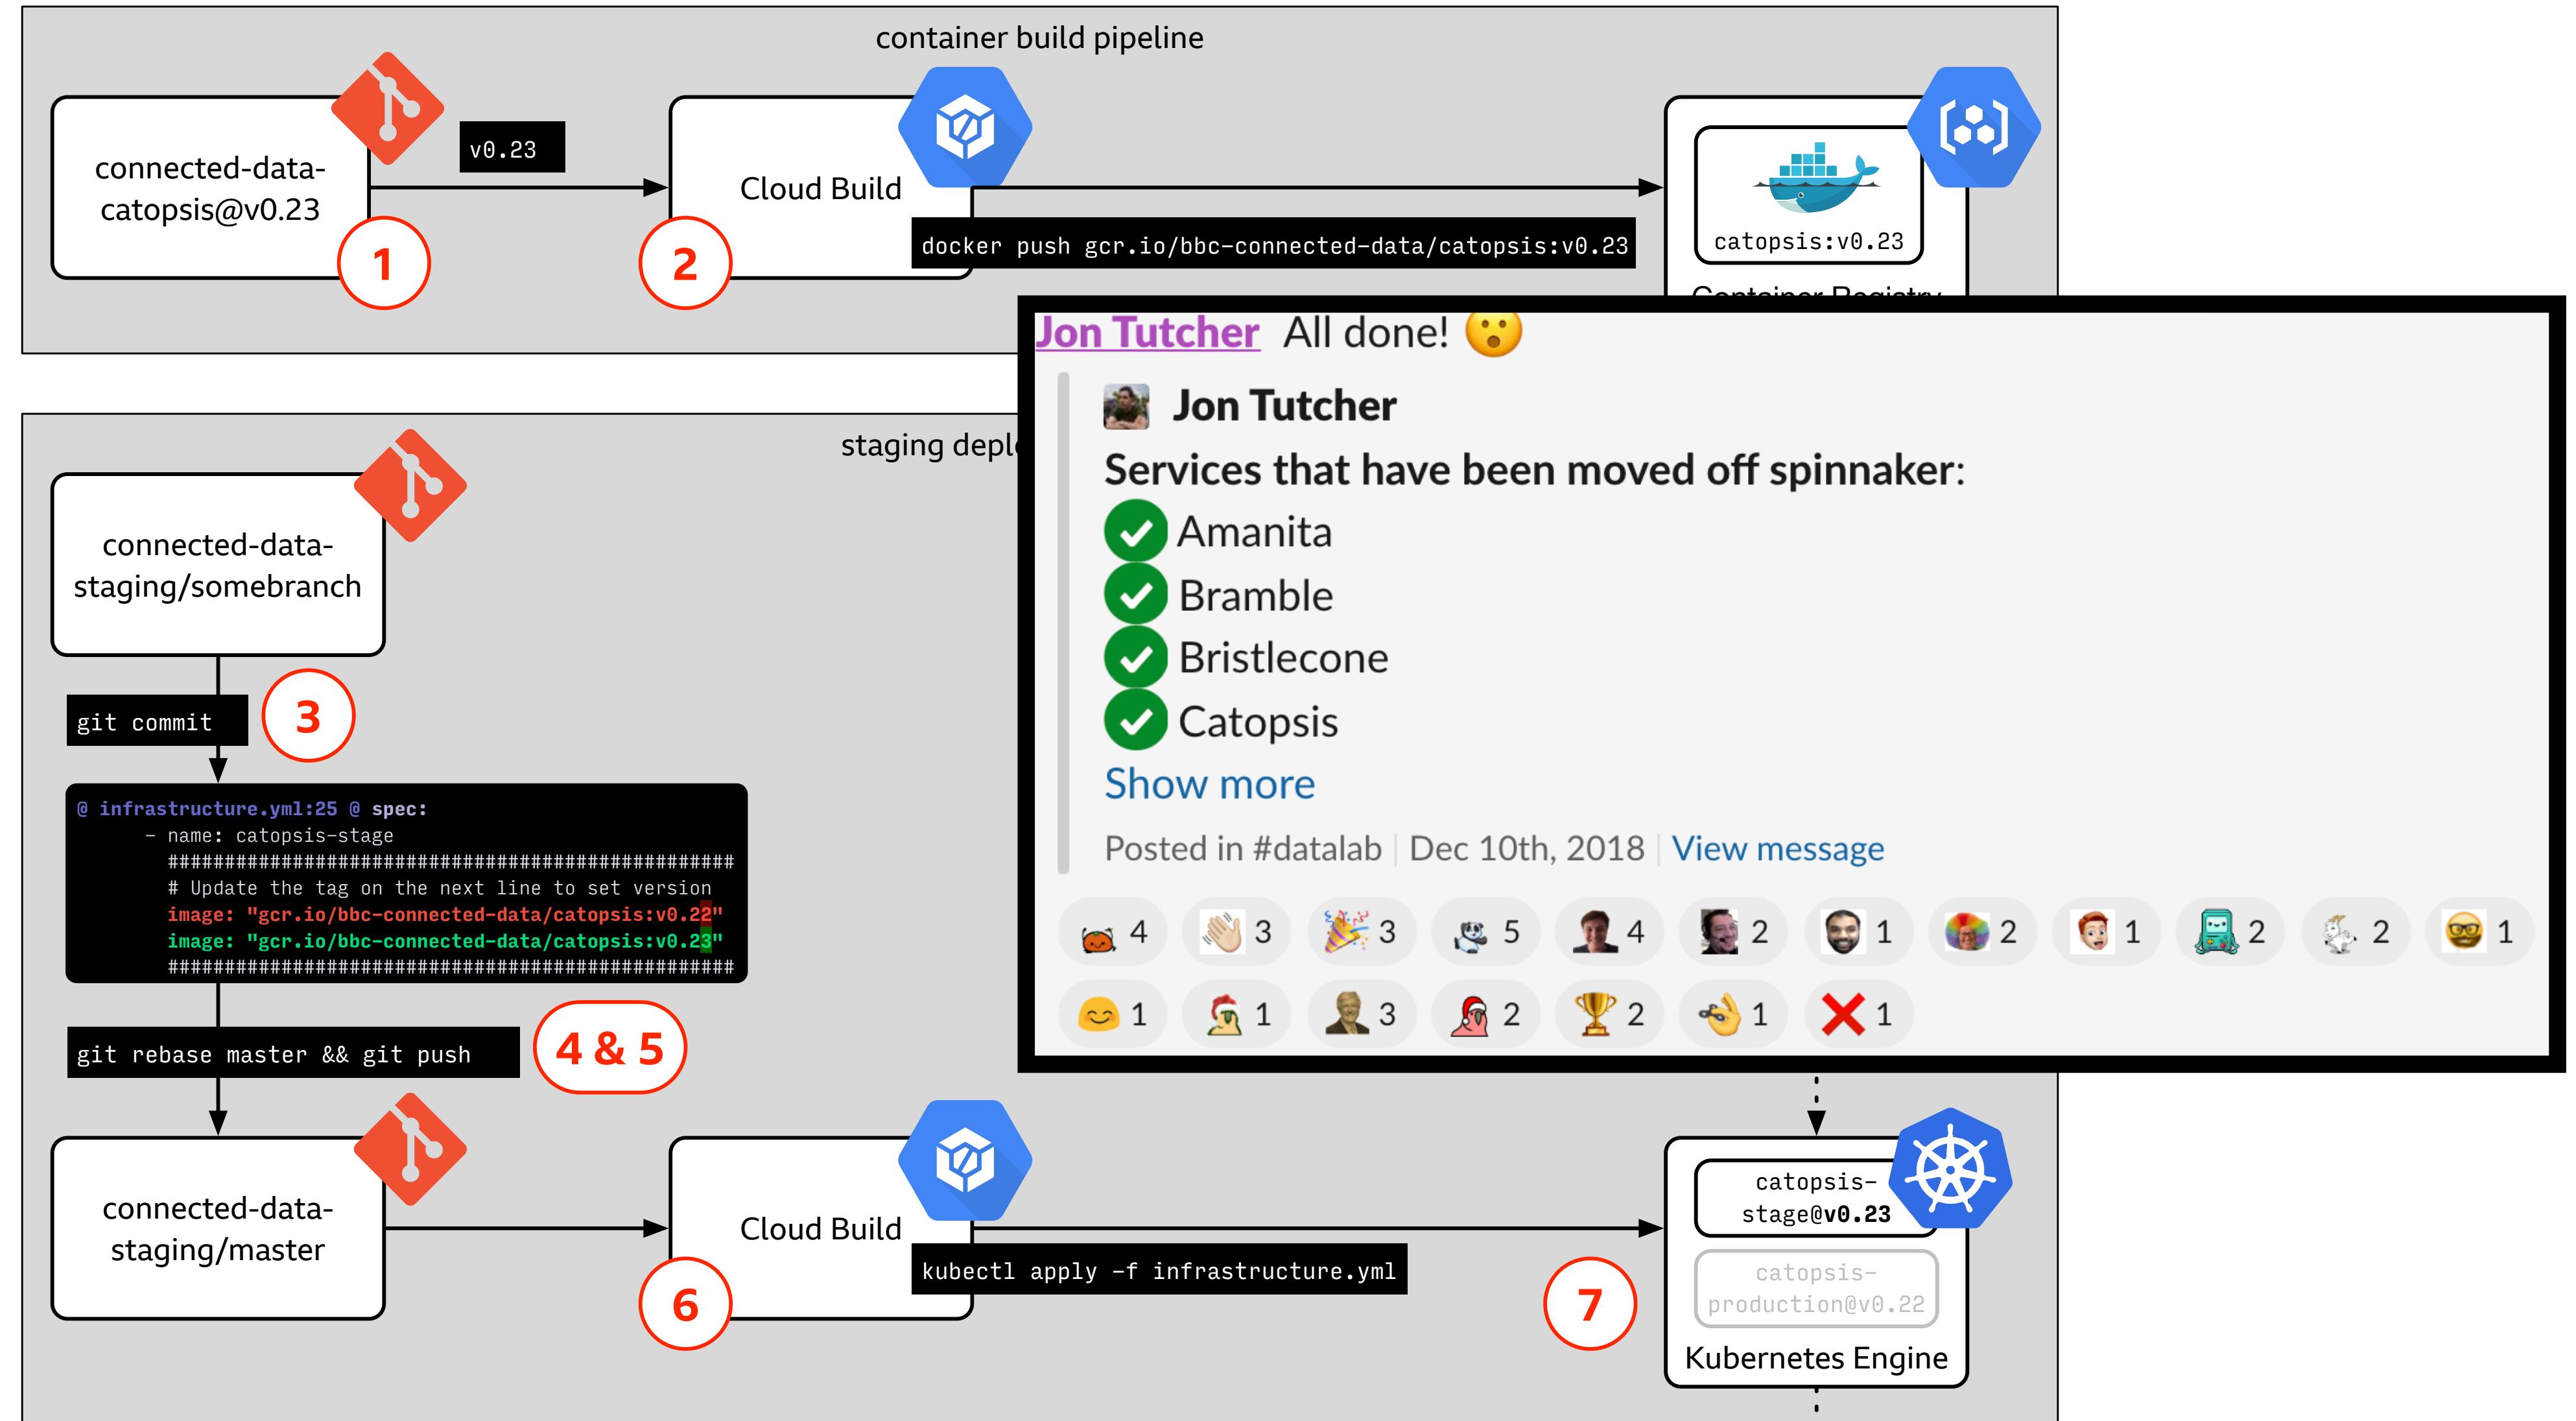
Created 2017-12-15 07:41:42 PST
In MY-K8S-ACCOUNT
Namespace default
Kind Service
YAML [Show YAML](#)
Kube UI [datalabservice-live](#)

Server Groups datalabservice-production-v006
datalabservice-production-v005

Service Type LoadBalancer

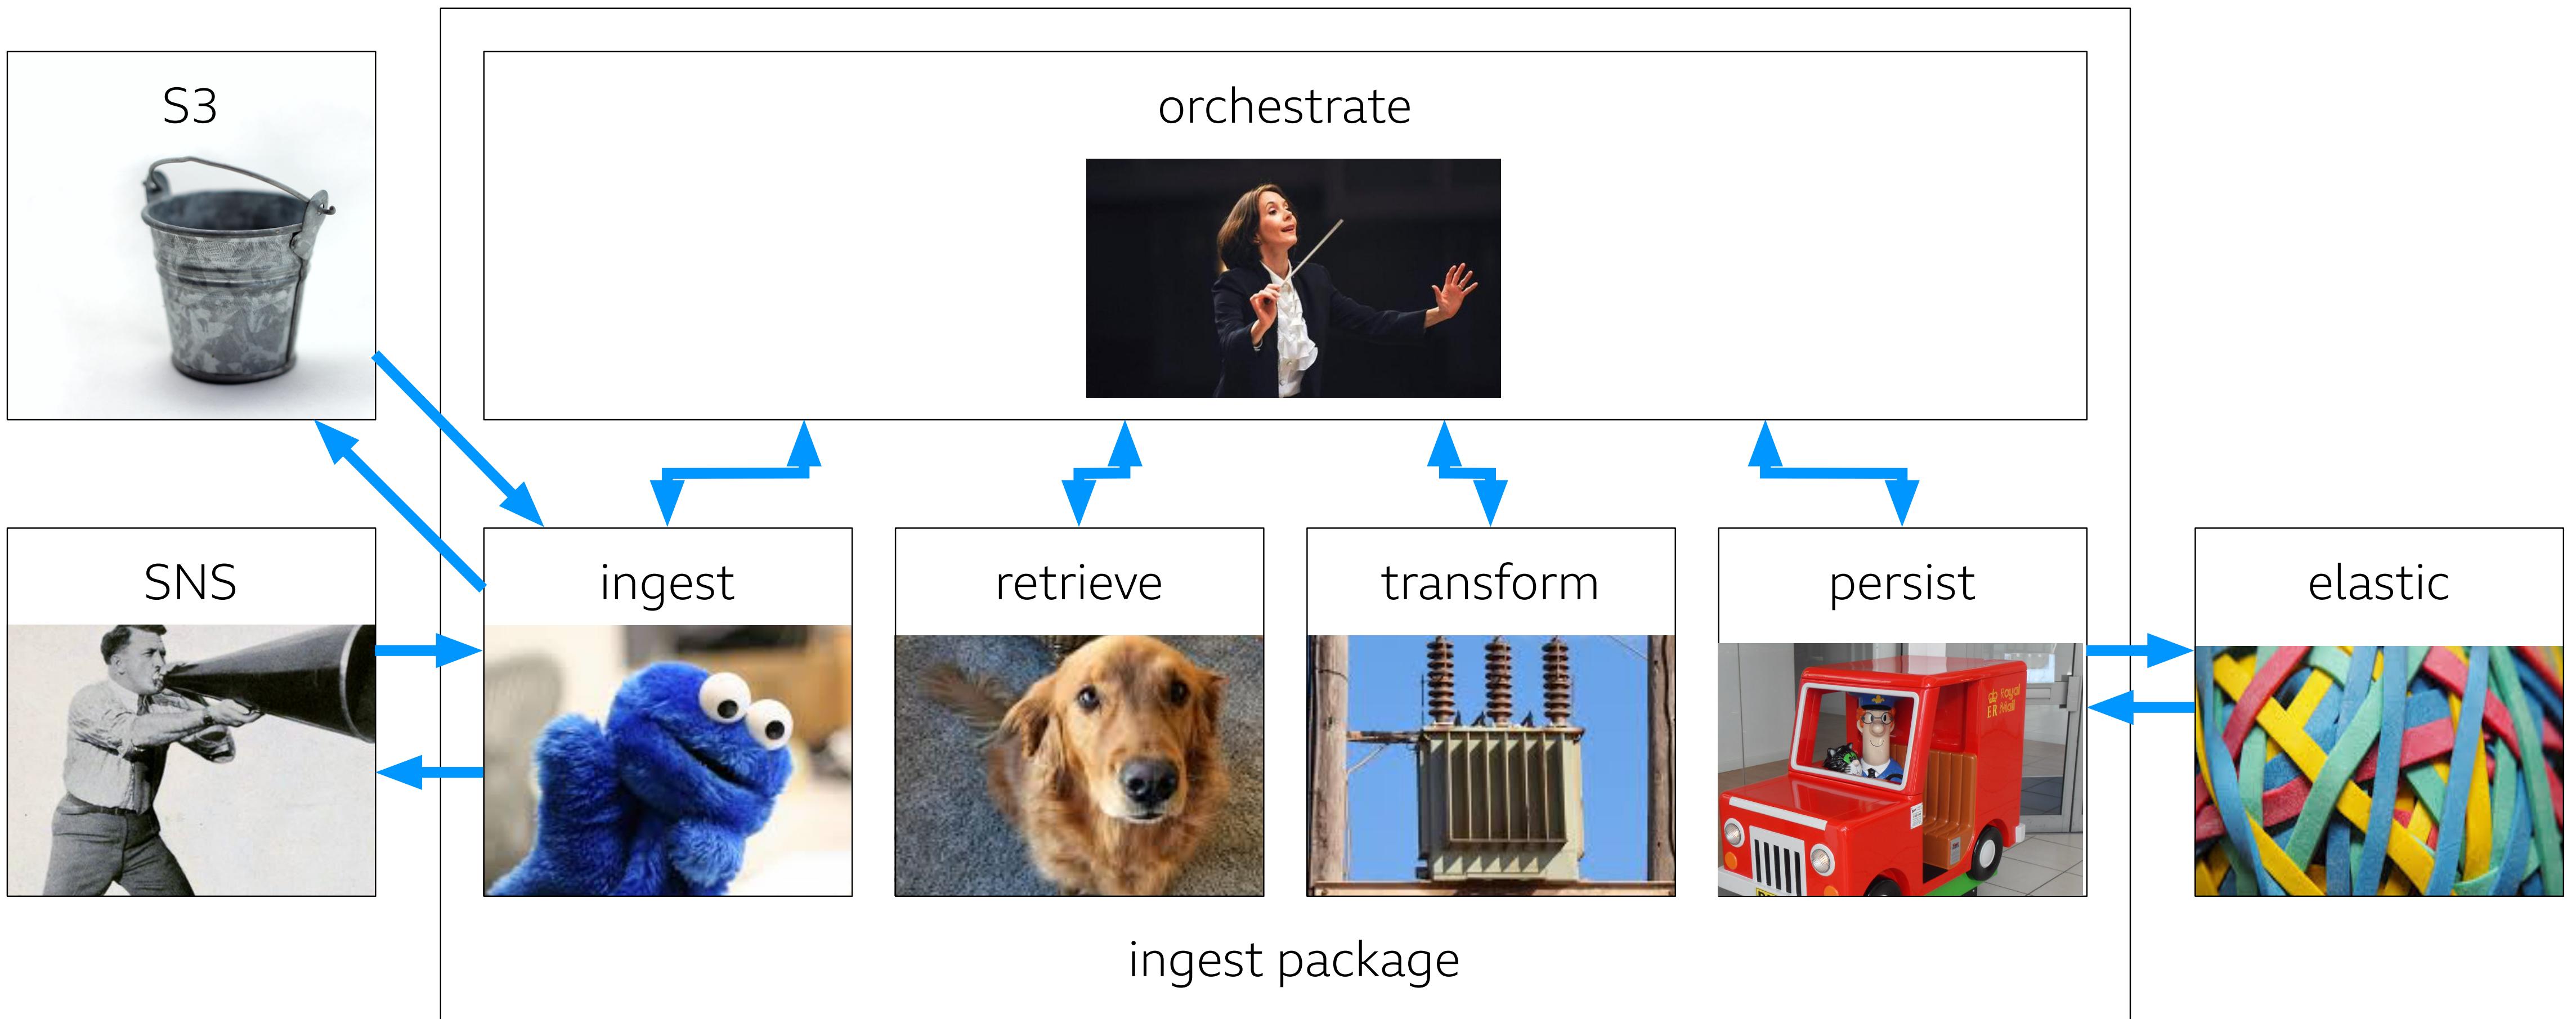
Session Affinity None

Internal DNS Name datalabservice-live.default.svc.cluster.local



OVERALL TEAM EFFECTS

- Pace dropped
- Low confidence in our code
- Data science dev slowed
- Bugs compounded
- Team morale dropped (until we started fixing!)



SEARCH



CONSOLE

- Change Kernel...
- Clear Console Cells
- Close and Shutdown...
- Insert Line Break
- Interrupt Kernel
- New Console
- Restart Kernel...
- Run Cell (forced)
- Run Cell (unforced)
- Show All Kernel Activity

FILE OPERATIONS

✓ Autosave Documents

- Close All
- Close Other Tabs
- Close Python File ^ Q
- Close Tabs to Right
- New View for Python File
- Open From Path...
- Reload Python File from Disk
- Revert Python File to Checkpoint
- Save Python File⌘ S
- Save Python File As... ⌘ ⌥ S

HELP

- JupyterLab Reference
- Launch Classic Notebook
- Markdown Reference

Launcher locustfile.py

```
1 import json
2
3 import numpy as np
4
5 import gevent.pool
6 from locust import TaskSet, Locust, task, HttpLocust
7
8 class HitApiEndpointTasks(TaskSet):
9
10     """
11     Define a set of tasks to run *per user/locust*.
12
13     Tasks
14         genres: a locust task that concurrently requests a number of videos, `LIMIT`,
15             for a number of genres, `N_GENRE_REQUESTS`. Genres are randomly sampled.
16     """
17
18     @task
19     def predict_flask(self):
20         LENGTH_OF_VECTOR=50
21         NUMBER_OF_VECTORS=2
22
23         vectors = []
24         for vector in range(NUMBER_OF_VECTORS):
25             topic_percentages = np.random.random(LENGTH_OF_VECTOR)
26             #the numbers in the vector should sum to 1
27             topic_percentages /= topic_percentages.sum()
28             vectors.append(list(topic_percentages))
29
30
31         locust_response = self.client.post('/predict', verify=False, data=str(vectors))
32         print("Response status code:", locust_response.status_code)
33         print("Response content:", locust_response.text)
34
35
36 class HitTfServeEndpointTasks(TaskSet):
37     @task
38     def predict_tfserving(self):
39         LENGTH_OF_VECTOR=75
```

TECHNOLOGY CHOICES

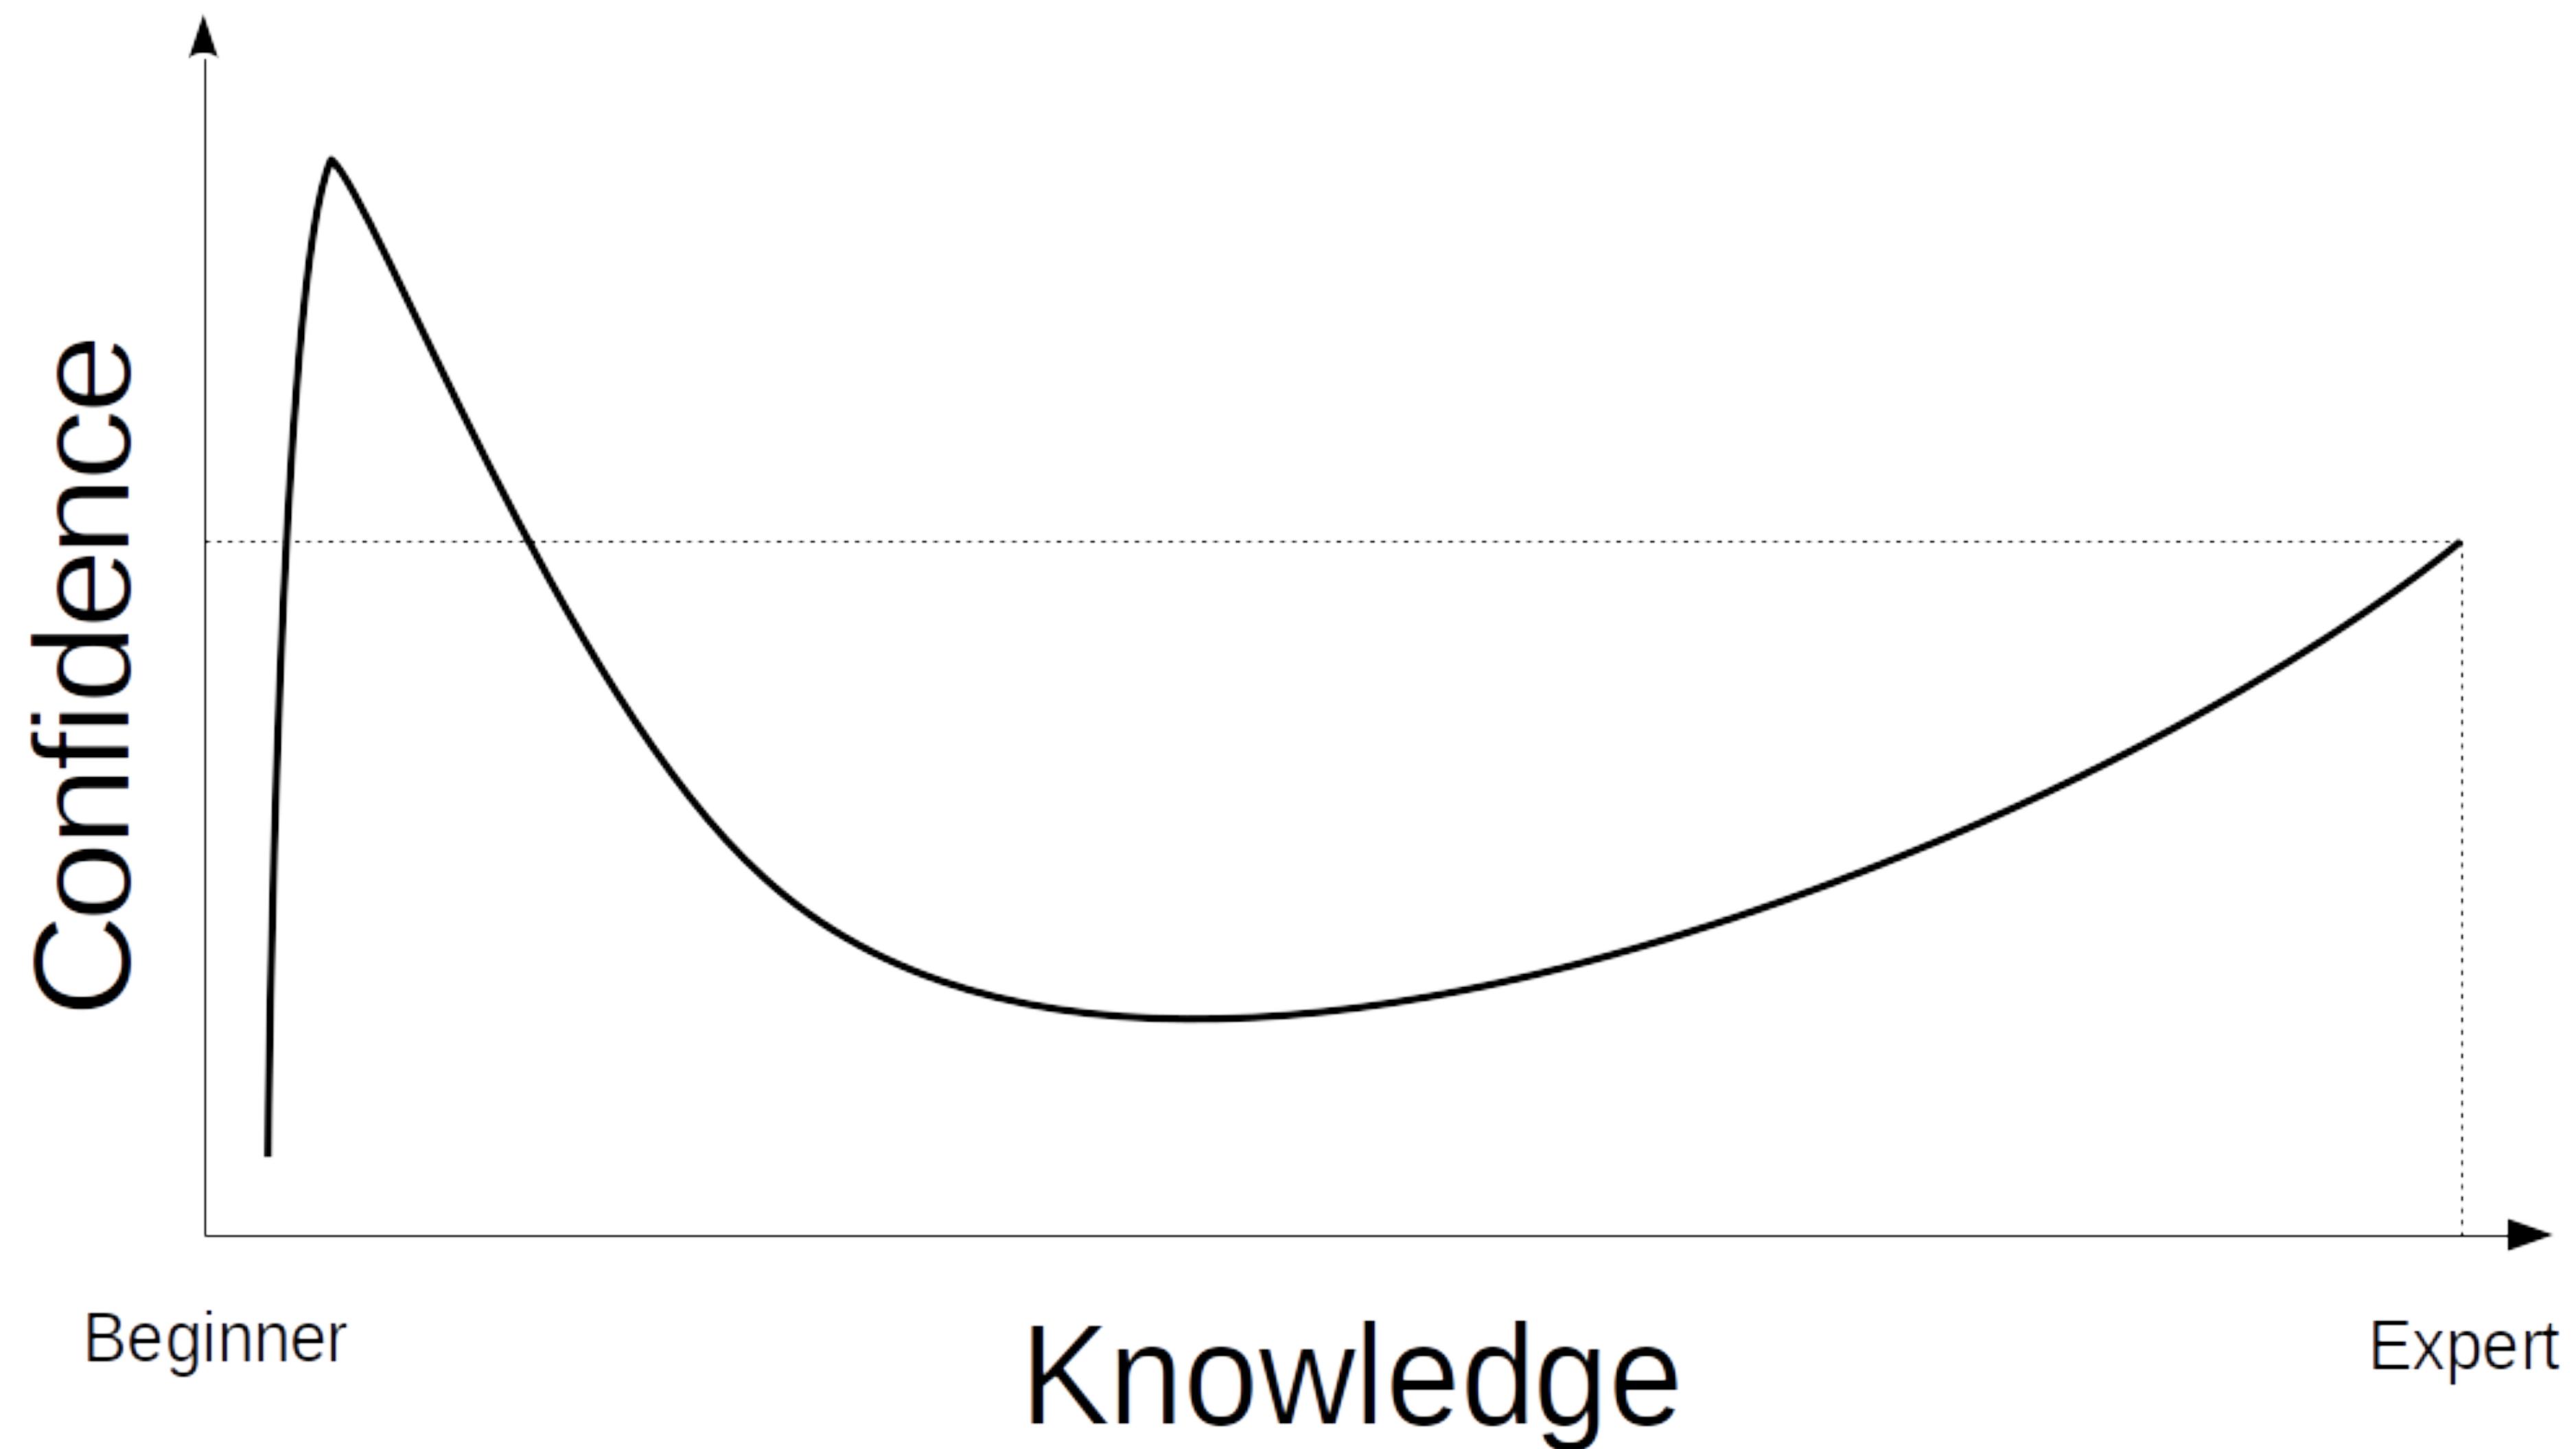
OTHER GOOD IDEAS

- Elasticsearch
- Managed Logging (mostly)
- Managed Training (Google ML Engine) (mostly)
- Committing to Tensorflow (for now)

TECHNOLOGY CHOICES

LESSONS LEARNT

- Decision making in new teams
- Over-engineering is easier than doing research
- Selection bias in press / meetups
- Python = no hassle
- Kubernetes = a keeper (for larger projects)

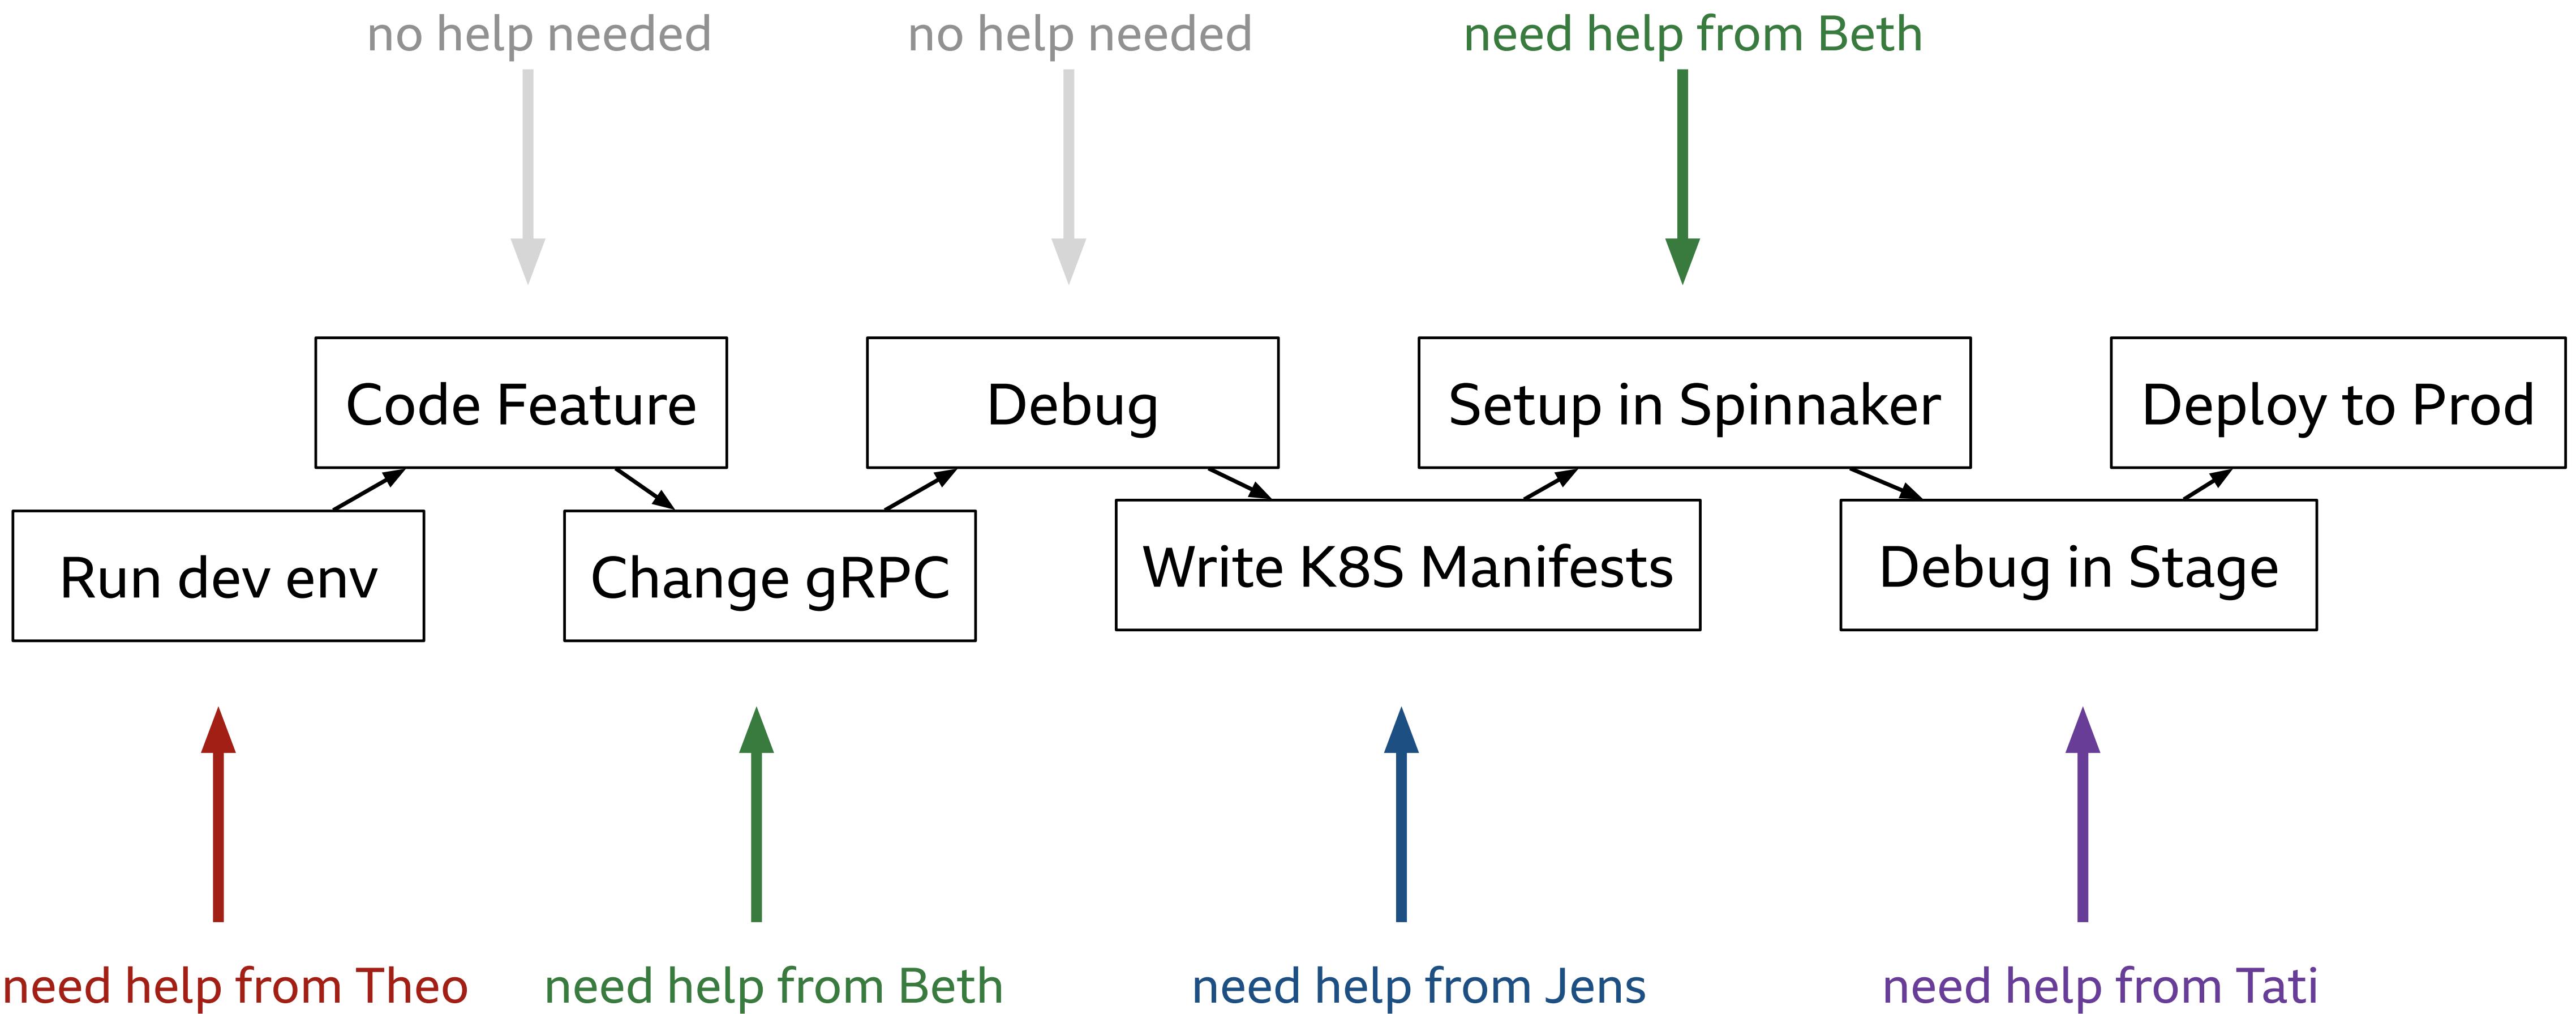


"USE BORING TECHNOLOGY"

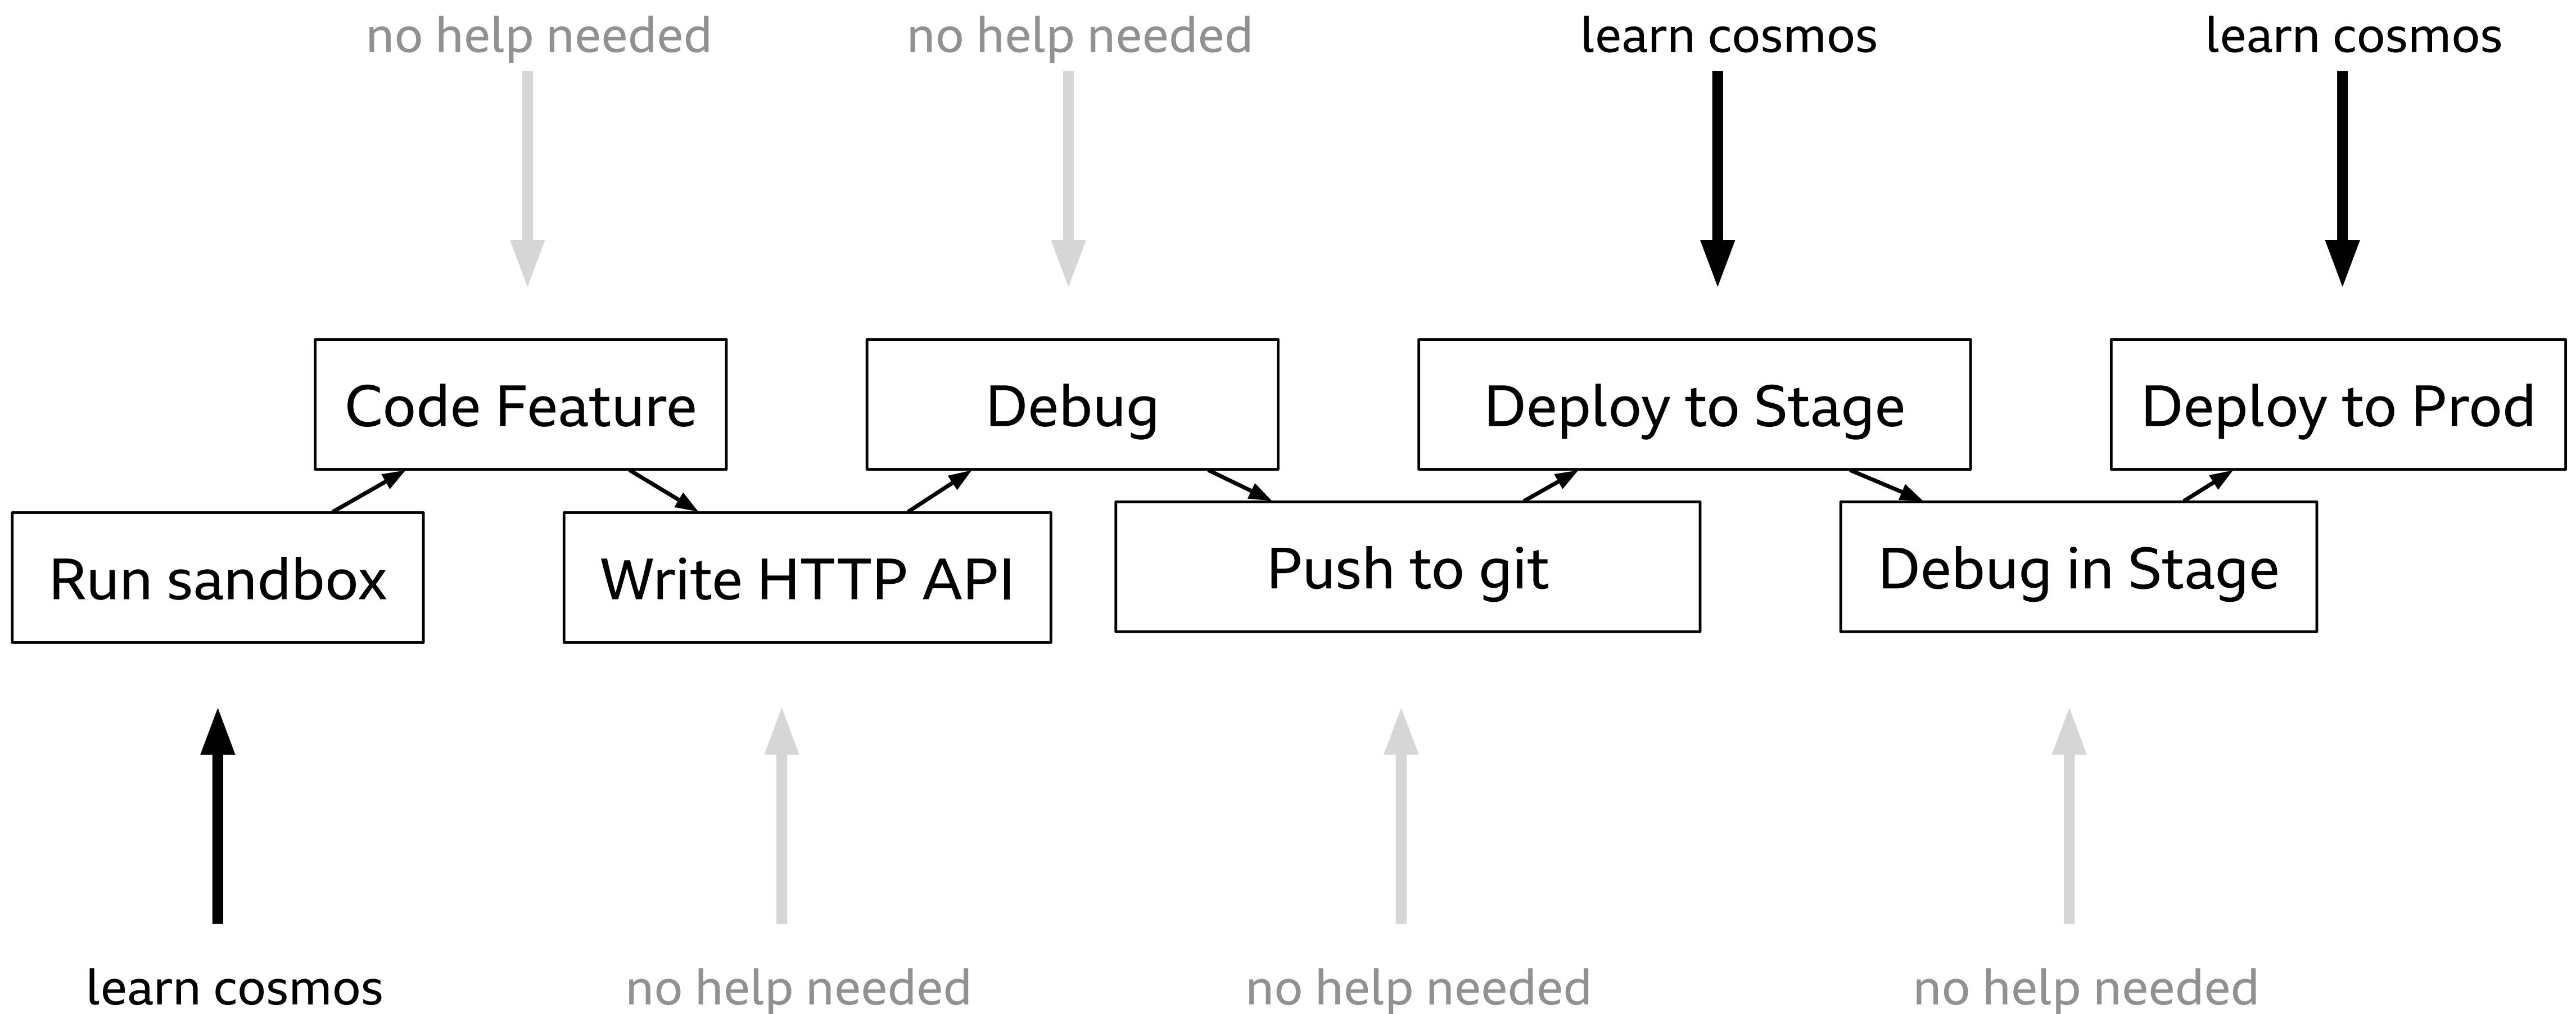
"The grim paradox of this law of software is that you should probably be using the tool that you hate the most. You hate it because you know the most about it."

- Dan McKinley, <http://boringtechnology.club/>

ANATOMY OF TASK BLOCKING



ANATOMY OF TASK BLOCKING



*cosmos = BBC's cloud deployment platform

```

class Technology:
    def __init__(self, name, maturity, familiarity, support, maintenance_cost, benefit):
        self.name = name
        self.maturity = maturity
        self.support = support
        self.familiarity = familiarity
        self.maintenance_cost = maintenance_cost
        self.benefit = benefit

    @property
    def pace_cost(self):
        risk = (1 - self.maturity) + (1 - self.support)
        return risk * (1 - self.familiarity)

    @property
    def total_cost(self):
        benefits = self.benefit
        risks = self.pace_cost + self.maintenance_cost
        return max(risks - benefits, 0)

technologies = [
    Technology("Spinnaker", maturity=0.1, familiarity=0.2, support=0.4, maintenance_cost=0.7, benefit=0.5),
    Technology("Postgres", maturity=1.0, familiarity=0.8, support=1.0, maintenance_cost=0.5, benefit=0.7),
    Technology("Hosted SQL", maturity=0.7, familiarity=0.5, support=0.7, maintenance_cost=0.2, benefit=0.7),
    Technology("Airflow", maturity=0.2, familiarity=0.5, support=0.4, maintenance_cost=0.4, benefit=0.6)
]

# Model cost of adoption
for tech in technologies:
    print(f"{tech.name}: pace cost: {tech.pace_cost:.2f}, total cost: {tech.total_cost:.2f}")

```

```

Spinnaker: pace cost: 1.20, total cost: 1.40
Postgres: pace cost: 0.00, total cost: 0.00
Hosted SQL: pace cost: 0.30, total cost: 0.00
Airflow: pace cost: 0.70, total cost: 0.50

```

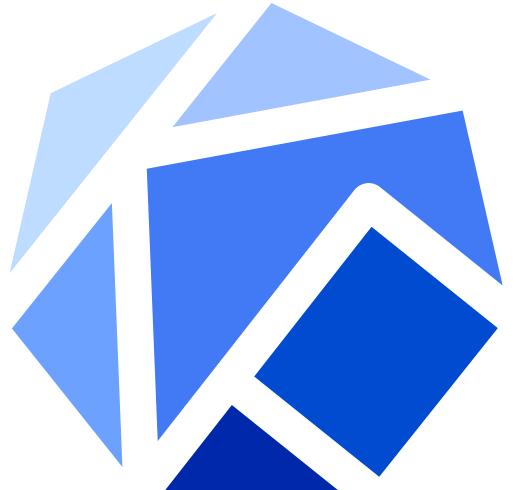
BBC+

THE NEXT CHALLENGE

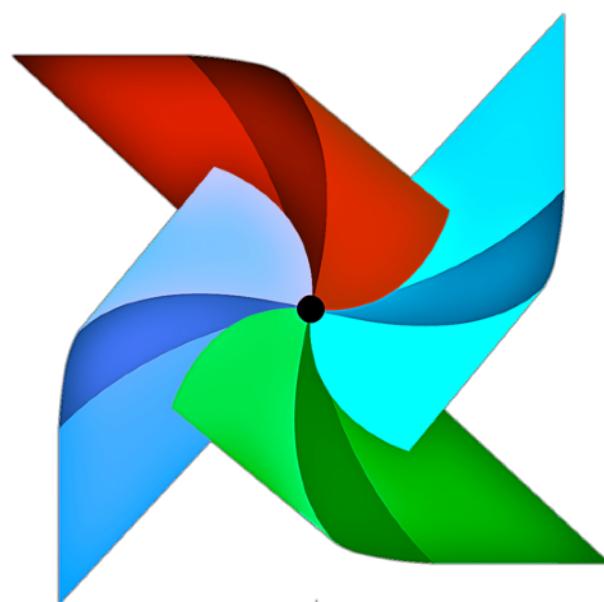
"How can we make model exploration and creation as automated as possible, whilst tracking provenance of data and code"?

WHAT NEXT?

THE KNEEJERK REACTION



Kubeflow



TensorFlow
Extended



SELDON



ML WORKFLOW TOOLS

WHAT'S NEXT?

```
new_techs = [Technology("Luigi", 0.2, 0.1, 0.3, 0.7, 0.7),  
             Technology("MLflow", 0.3, 0.2, 0.3, 0.7, 0.7),  
             Technology("Dask", 0.4, 0.3, 0.5, 0.6, 0.7),  
             Technology("Kafka", 0.6, 0.7, 0.7, 0.7, 0.7),  
             Technology("Beam", 0.4, 0.5, 0.4, 0.7, 0.8),  
             Technology("Jenkins", 0.9, 0.7, 0.9, 0.2, 0.6)]
```

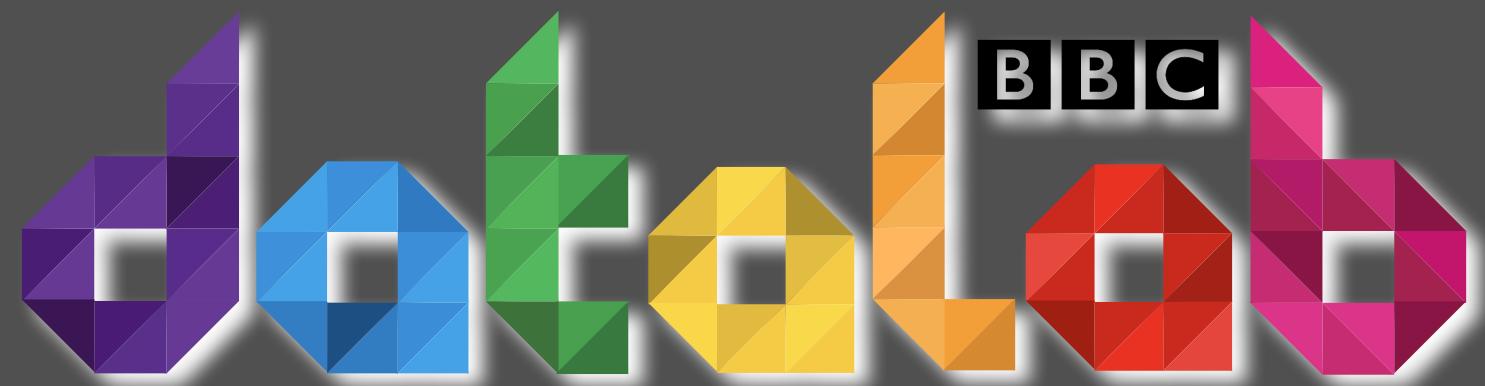
```
Luigi: pace cost: 1.35, total cost: 1.35  
MLflow: pace cost: 1.12, total cost: 1.12  
Dask: pace cost: 0.77, total cost: 0.67  
Kafka: pace cost: 0.21, total cost: 0.21  
Beam: pace cost: 0.60, total cost: 0.50  
Jenkins: pace cost: 0.06, total cost: 0.00
```

FINAL THOUGHTS

- Fit your problem to existing tech (if poss)
- Avoid sunk cost fallacy
- Experiment, but one-at-a-time
- What's right for Google isn't right for you

THANKS!

@jontutcher



Come and work with us!

<https://findouthow.datalab.rocks/>